

HM - Interview

HM Interview

[Greeting](#)

[Ques to ask in the end](#)

[Introduction \(45-60 secs\)](#)

[Why Solutions Architect?](#)

[SF - influenced a customer architecture \(1-2 mins\)](#)

[Apple - Spark and Data Lake / Lakehouse](#)

[Data Governance Exp](#)

[Presales Experience?](#)

[IRhythm - \(buzz words - sales POC, 2 weeks workshop\)](#)

[BayRen RFP](#)

[AWS MAP Funding](#)

[Spark Job Optimization](#)

[Customer Disagrees \(refined layer\)](#)

[Calculated Risk -\(Speed / best practices first\)](#)

[Frustrated/impatient customer \(Resetting expectations\)](#)

[Failed Architecture: over-engineered \(separate jobs/per DS\)](#)

[Over-commitment \(all dashboards\)](#)

[Above and beyond the duty \(PM hat\)](#)

[Coaching Customer team.\(Influencing team\)](#)

[Simplifying for customer - Pipeline self-service](#)

[Machine Learning confidence?](#)

[Comfort with databricks stack?](#)

[Top Strength](#)

[Weakness?](#)

[Future: see yourself in 5 years from now?](#)

[What does day to day look like?](#)

Greeting

“Hi XYZ, it’s really nice to meet you.

Ques to ask in the end

1. “What does success look like in the first 3–6 months for someone stepping into this SA role—both technically and from a customer-impact standpoint?”

2. “Is there anything in my background that you’d like me to elaborate on, or any concerns I can help address?”

Introduction (45-60 secs)

I’m a Solutions Architect focused on Data & Analytics at Slalom. I’ve been **leading** the **design** and delivery of **distributed** data platforms using **Spark**, **Iceberg**, and **cloud-native** services.

And I work directly with **business/leadership stakeholders** and engineering **teams of clients** across both **pre-sales** and **delivery tracks**.

A **big part of my role** is **guiding customers** on **architecture, governance, performance, and cost trade-offs**, and **building/running POCs** in their environment.

I’m **passionate** about helping customers solve **real business problems** by strengthening their data engineering capabilities **so they can** move faster, **operate efficiently**, and get more value from their data.

I’m **particularly inclined towards Databricks** because of its lakehouse platform, **its open-source leadership**, and the customer-facing SA role aligns closely with my experience.

That’s a bit about me, and I am happy to discuss more about my experience.

Why Solutions Architect?

I’ve **always gravitated toward roles** where I can **connect** deep **technical architecture** with **real business impact**.

“I’m passionate about helping customers solve real business problems by strengthening their data engineering capabilities so they can move faster, operate efficiently, and get more value from their data.”

This SA role is exactly at the intersection of what I enjoy and what I’m good at — customer-facing architecture, data platform modernization, and translating complex technical decisions into business value.

I’ve spent the last several years helping customers design and modernize data platforms — guiding them from legacy systems to scalable, governed, analytics- and AI-ready architectures. What excites me about this role is the opportunity to do that at *Microsoft’s* scale, using *Azure’s* ecosystem — Databricks, Fabric, Purview — and to influence not just designs, but adoption and long-term success.

Why Databricks?

Leaders and inventors in the industry of many data engineering solutions - lakehouse, unity catalog,

I'm drawn to Databricks because **it consistently solves the advanced data engineering problems** I see customers struggling with. The Lakehouse architecture and Unity Catalog address these challenges in a way that's both technically strong and straightforward to operate.

I also value Databricks' open-format, vendor-neutral approach. Customers want flexibility, not lock-in, and Databricks enables that while still giving them a unified platform for data, analytics, and AI.

For me, joining Databricks means being in a place where technical depth directly translates to customer impact—helping teams move faster, reduce complexity, and get more value from their data.

SF - influenced a customer architecture (1-2 mins)

One of the most **impactful architectures I influenced** was for a large CRM company whose analytics team was **constrained by fragmented data, inconsistent data quality, and recurring compliance issues**. Their business teams **needed a more scalable and reliable way** to report on compensation data

I started with a **discovery phase where I interviewed around 15 stakeholders** across engineering, analytics, and compliance. I reviewed their end-to-end data flows and found that roughly a third of the upstream data feeding their analytics environment wasn't SOX compliant. **I consolidated these and many other findings and presented, data-driven assessment to their leadership.**

I recommended two future-state options — S3 + Snowflake, and an S3-based lakehouse — and collaborated with their enterprise architecture team to make sure the design fit the company's overall platform direction. I presented the architectural options to the leadership team and C-level execs and 3v

This was very well received by the customer and their leadership.

I also helped define a **phased implementation plan** that delivered **quick wins**. With the new architecture and the automation we introduced across ingestion, validation, and reporting workflows, **the team reduced manual effort significantly and accelerated time-to-insight by nearly 80%**.

Apple - Spark and Data Lake / Lakehouse

I worked with an **AI/ML org of a top technology company** that builds **intelligent features for their iPhone**, and was supported by **6,000 global annotators**. Their operations were fully data-driven. Metrics like **Productivity, quality**, and compliance are **continuously tracked through dashboards and performance analytics**.

Their existing **data ingestion process was not scalable**, another challenge was that the **upstream data provider team** was not consistently **delivering data in the structure or cadence the data scientists needed for analysis**.

I started by **proposing a fast MVP using 2–3 important datasets to show quick value** and set up a foundation for scale. I designed a **lakehouse** architecture using **AWS Glue + PySpark, Iceberg tables, Airflow orchestration, and a medallion layout (raw → refined → curated)**.

I used Great Expectations for data-quality checks, and I worked closely with data scientists to **curate datasets at the right granularity** for their models.

I demonstrated how Apache Iceberg enables **time travel, incremental refresh, schema evolution**, while giving us stronger **observability**, and **traceability** through its metadata and versioning.

The self-service analytics pipeline we built let the business users explore and onboard datasets independently, and **it drove a 90% improvement in data accuracy**.

This led the customer to adopt this architecture for all datasets **and expand the engagement with a larger team**. It was the same engagement where the a senior stakeholder awarded me the “customer impact ward”

Data Governance Exp

I've worked with several customers on data governance, **and one thing I've learned is that priorities vary.** Some teams **care most about discovery and cataloging, while others focus on strict access controls or lineage.** But regardless of the priorities, **the foundational elements—cataloging, classification, ownership, access control, and quality—**need to be in place for any architecture to work.

In one customer engagement, I was initially brought in to help with data engineering, but as trust grew, they depended on me to shape their overall governance approach and ensure it aligned with their architecture roadmap.

I **reviewed their existing access controls** and redesigned them using Snowflake RBAC, introducing clear role hierarchies, least-privilege principles, and standardized permission patterns.

A major gap was the ambiguity around data ownership and stewardship. **I facilitated alignment across analytics, engineering, and business teams and created a simple RACI** that clarified who owns datasets, who publishes them, and who approves access. This reduced friction significantly and higher trust in the data.

Presales Experience?

I've done quite a bit of pre-sales work, especially in **shaping architecture options, running short POCs, and contributing to RFPs.** My focus is usually on understanding the customer's pain points quickly, proposing a practical approach, and then de-risking the decision with something tangible.

For example, **with a new health-tech client, I led a four-week sales POC** to help them evaluate a new data foundation. I scoped the use cases, designed the architecture, and demoed capabilities like automated cataloging, PII detection, and data-quality rules. **That early value accelerated their** time-to-insight by about 80% on the first few datasets and helped our team win the follow-on work.

Beyond that, **I regularly support RFPs—**writing the technical architecture sections, solution assumptions, and delivery model—**and I've also secured AWS MAP funding to reduce POC cost barriers for clients.**

Overall, my pre-sales experience is hands-on: understanding the problem, shaping the right solution, validating it quickly, and helping the sales team move the opportunity forward.

IRhythm - (buzz words - sales POC, 2 weeks workshop)

I was engaged early with a new health-tech client — **a net-new logo for Slalom** — whose data services organization needed to **deliver faster, more reliable insights from large volumes of device-generated data**. They were in the middle of a **Redshift-to-Snowflake migration**, but **lacked a consistent ETL/ELT framework**. Data ingestion was driven by new Go-based extractors, and data movement, transformation, and validation were slow, manual, and difficult to govern.

After assessing the environment and aligning with key stakeholders, I **proposed a focused four-week “test-and-learn” proof of concept** to demonstrate what a modernized data foundation could look like. The POC prioritized three capabilities:

- Automated data discovery and cataloging
- PII-aware metadata tagging and privacy compliance
- Data-quality monitoring across representative datasets

I designed and demoed the solution using **AWS Glue** for cataloging and schema tracking and **Amazon Macie** for PII detection, **integrated with quality rules and remediation workflows**. I presented the target architecture to both executive and engineering audiences, addressing their questions around privacy, security, cost trade-offs, and platform alternatives.

At the end of the four weeks, we delivered a functioning end-to-end data flow with a metadata- and PII-aware catalog dashboard. This accelerated **speed-to-insight by ~80%** across the initial 2–3 datasets and gave the client the confidence to move into the next phase of execution with Slalom as their strategic partner.

BayRen RFP

“One of my key pre-sales engagements was with BayREN, a regional agency that runs six different energy-efficiency programs. Each program had its own database and reporting process, so leadership had no unified way to see performance across the region. The goal of the RFP was to design a data aggregation and reporting platform that could automatically pull and standardize data from all six programs into one reporting layer.

The main challenges were fragmented systems, inconsistent schemas, and completely manual reporting. They needed a common data model, automated pipelines, and a governed environment — but this was still at the proposal stage, so the focus was on shaping a practical, scalable approach.

My role was as the technical solution lead for the RFP. I worked with our practice leads, architects, and the client's evaluation team to define requirements and shape the architecture. I wrote the technical sections of the proposal, including the end-to-end data platform design on AWS using S3, Glue, Redshift, and Lake Formation, along with assumptions, constraints, and a clear delivery approach.

I also helped draft the SOW, RACI, and staffing model, and joined the orals presentation to walk both business and technical stakeholders through the proposed design. This work helped Slalom advance to the final selection stage and contributed to about \$2–3M in pipeline for the Data & Analytics practice.”

AWS MAP Funding

“We had a healthcare client who wanted to test a new data architecture but **didn't have the budget for a proof of concept**. We didn't want to lose momentum, so I was asked to explore funding options. I worked with our AWS alliance team and looked at the MAP program, which can fund POCs that drive future cloud adoption.

I drafted a short proposal outlining the use case, the AWS services involved, and how it would grow their long-term AWS consumption. The proposal was approved, and AWS fully funded the POC. I then led a small cross-practice group to build the demo using AWS analytics services in the client's environment.

The client got a working POC at zero cost, it built a lot of trust, and it opened the door for a larger engagement. It also strengthened our partnership with AWS for future opportunities.”

Spark Job Optimization

In one of my past projects, prior to Slalom, I had a Spark job on AWS EMR processing loan events from Kinesis. **The job was running on Spark 2.x and was consistently missing its 5-minute micro-batch window**, which caused batch queues, high scheduling delay, and ultimately data loss that affected a large set of customers. Upgrading to Spark 3 was the long-term plan, but we needed an immediate fix to stabilize production.

I started by profiling the job in the Spark UI — which wasn't even enabled due to firewall rules, so I worked with the networking team to expose it in production. From there, **it was clear the main issues were data skew and inefficient joins**. One dataset had a heavy skew on a few key values, so certain partitions were 10–15x larger. I added targeted filtering before joins, applied salting to reduce skew, and optimized the aggregation logic.

On the infrastructure side, I recalculated the executor/cores configuration based on the actual workload instead of the default EMR settings, which improved parallelism and reduced shuffle overhead.

After these changes, the micro-batches consistently completed within the 5-minute window, processing time dropped significantly, and we stopped losing customer events.

This stabilized the workload and improved end-to-end processing efficiency by roughly 80%, allowing the system to meet the 5-minute SLA while the Spark 3 upgrade was being planned.

Customer Disagrees (refined layer)

*Themes - understand constraint, **co-design pragmatic path**, trade-offs*

When a customer disagrees with a recommendation, I **step back and understand the constraint — it could be cost, ownership, or complexity**. I keep it data-driven and look for a solution that aligns with their priorities.

One of my customer's Data Science & Analytics team was working off nearly a million upstream transactions a day, but the raw data had inconsistencies that impacted feature engineering and model accuracy. My architecture recommendation had a refined layer for cleansing, but they pushed back due to long-term data-engineering costs. **Their concern was valid — and waiting months for the upstream team to fix the data wasn't an option because their modeling roadmap would stall.**

So, I proposed a **lighter-weight approach**:

Build a YAML-configurable cleansing layer so they get reliable, analysis-ready data immediately with minimal engineering.

Automatically generate structured outputs of all cleansing rules and pass them to the upstream data owners so the logic can be absorbed into their pipelines over time.**

This gave the Data Science team clean, consistent data quickly without taking on permanent maintenance. Within a quarter, data-quality issues dropped by ~95%, and their team could focus on feature development and modeling instead of fixing data.

"I focus on **aligning with the customer's constraints and co-designing the most pragmatic path** forward, rather than defending a single solution."

Calculated Risk -(Speed / best practices first)

Themes - architecture hardening later, value fast, Build trust through small, predictable value cycles

Give me an example of a calculated risk that you have taken where speed was critical. What was the situation and how did you handle it?

An example of a calculated risk was on a project where I joined right after the design phase—and the consultant who owned the pipeline architecture had just left. The customer's Senior Director urgently **needed a high-priority dataset** to unblock both their BI reporting and machine-learning modeling, but the infrastructure provisioning—owned by a separate team—was delayed by several weeks.

Speed mattered because their forecasting team was stuck with no data. **Instead of waiting for full CI/CD, networking, and IAM patterns to be finalized**, I proposed delivering a lightweight **MVP ingestion-to-curation pipeline directly in the AWS console using Glue jobs, S3 staging, and parameterized transformations**. **It wasn't the long-term architecture, but it allowed us to deliver the critical dataset in under 5 days—versus the 3–4 week delay the customer was expecting.**

This was a risk, because the Engineering Manager **preferred a strict “best-practices-first” approach**. I aligned him by clearly laying out the two-track plan:

Track 1: Deliver the urgent dataset so ML and BI teams could begin exploratory modeling, which unblocked two downstream workstreams and ~8 analysts.

Track 2: In parallel, harden the pipeline into a production-grade framework with CI/CD, modular jobs, and cost-optimized orchestration.

The outcome: the customer started generating early features and dashboards immediately, **leadership saw value within the first sprint**, and the Engineering Manager supported the roadmap because he saw how the MVP accelerated time-to-insight without compromising the long-term architecture.

Frustrated/impatient customer (Resetting expectations)

Most of us at one time or another have felt frustrated or impatient when dealing with customers. Can you tell us about a time when you felt this way and how you dealt with it? When do you think it's appropriate to push back or say no to an unreasonable customer request?

Answer:

In one meeting with a customer, a senior director said, 'We've been waiting months — **can't you just pull the data manually and patch the dashboards?**' I could feel their frustration —

and honestly, I felt it too. But **I knew saying yes would create permanent technical debt and destroy trust later.**

First, I acknowledged the frustration, by telling them

‘I agree this has been painful. You’re not getting what you need to support a 6,000-person global annotation workforce — the impact is real.’

Then I used data to reset expectations. I showed that:

- **Upstream changes were breaking schemas weekly.**
- **Without a governance layer, we couldn’t guarantee any dashboard’s accuracy.**

I explained that patching would only give them another failure in 2–3 weeks.

So I proposed a **fast, low-risk MVP using Iceberg — showing how time travel, schema evolution, and metadata tracking would fix quality issues at their root.**

I push back when:

- The request introduces risk
- It bypasses governance,
- Or it gives short-term relief but long-term pain.

I do not say no bluntly — I explain the ‘why’ with numbers and offer a better path. Customers appreciate that level of honesty.

Failed Architecture: over-engineered (separate jobs/per DS)

Themes: Ownership, simplification, redesign, operational efficiency

“In one engagement, **I designed an ingestion framework on AWS Glue where each data source had its own job. That pattern is usually helpful because it gives isolation for schema differences, SLAs, and independent deployments.**

But after we got further into the build, it became clear that this **customer’s workload was fairly stable and wasn’t expected to grow or evolve frequently.** They had about a dozen sources with simple transformations that rarely changed. In that context, the one-job-per-source pattern created unnecessary overhead — separate schedules, logs, and monitoring points — without giving proportional value.

I recognized that the design, while technically correct, wasn't the best operational fit for a small, steady-state workload. I took ownership of that and walked the customer through the trade-offs: more modularity versus more operational touchpoints. We aligned on a better approach — a metadata-driven, parameterized Glue job orchestrated through Step Functions. This gave them a single operational surface area, consistent retries and monitoring, and still allowed source-specific logic through parameters.

After implementing the revised architecture, maintenance became much simpler, and the customer appreciated that we optimized the design for their actual long-term needs rather than sticking to a pattern just because it's "best practice."

Over-commitment (all dashboards)

Themes: Scoping, prioritization, exec communication, delivering with quality

To try to meet the high expectations of our customers, we sometimes promise more than we can deliver. Tell me about a time when you overcommitted yourself or your company. How did you resolve the issue?

One time I overcommitted was with a customer who needed a full reporting suite **delivered within the quarter**. Based on the components already in place, it initially looked achievable, so I committed.

But as we got deeper into the dataset and business logic, the scope expanded—more metrics, more drill-downs, and stricter data-quality requirements. It became clear we couldn't deliver all 12 dashboards at the expected fidelity without risking quality.

I paused the team and did a rapid assessment of **effort vs. value across all use cases**. We identified the top 6 executive dashboards that directly supported the customer's quarterly leadership review and could be delivered with high confidence. The remaining items—mainly deeper operational drill-downs—were important but not time-critical.

I discussed this with the customer, walked them through the trade-offs using data (complexity, row-level QA effort, refresh SLAs), and proposed a phased plan that preserved quality while still meeting their most critical goals. They appreciated the transparency and the fact that the priorities were tied to measurable business outcomes.

In the end, we delivered the **core dashboards on time and with 99% data accuracy**, and we sequenced the remaining work into the next sprint cycle. The customer met their leadership commitments, and we maintained trust by turning a risky overcommitment into a clearly managed plan.

Above and beyond the duty (PM hat)

Leadership, structure, escalation management, customer obsession

Tell me about a time when you went above and beyond the call of duty for a customer.

Why did you take the action you did? What was the outcome?

Midway through an engagement, the customer's expectations shifted. What started as a purely technical role suddenly required heavy project management, tighter coordination, and much more proactive communication across their engineering and analytics teams. I realized that if we didn't adapt quickly, delivery velocity and trust would both suffer.

To stabilize the program, I stepped in immediately. I took over and led daily stand-ups, reorganized the backlog, and set up a clear delivery cadence so the customer had full visibility into progress and risks. This alone improved issue resolution time by about 30%.

At the same time, I partnered closely with our Slalom account and delivery leadership to reassess the skill profile needed. It was clear the team needed someone with stronger PM and cross-functional coordination experience, so I moved quickly to initiate a fast-track hiring and onboarding process for a consultant whose background matched the new scope.

I stayed deeply involved through the transition to ensure zero disruption—maintaining sprint commitments and keeping the customer fully informed.

The outcome: the customer saw on-track delivery for the next two sprints, stakeholder satisfaction scores improved, and we strengthened our long-term relationship by showing that we could adapt rapidly and put their success first.

Coaching Customer team.(Influencing team)

Give me an example of how you have changed the direction or view of a specific function/department and helped them embrace a new way of thinking? Why was a change needed?

A strong example is when I helped a customer's Data & Analytics team shift from ad-hoc development to a more **scalable, governed data architecture**. Their biggest pain point was that each analyst and engineer was creating their own schemas, views, and pipelines. This led to **fragmented logic, duplicated datasets across 4–5 domains, inconsistent KPIs, and growing technical debt** that made reporting slow and unreliable.

When I reviewed their workflows with the engineers, I showed how a **medallion architecture (Bronze → Silver → Gold)** would create a single, reusable curated layer and reduce duplication. I conducted **deep-dive sessions** to walk them through schema standardization, data-quality rules, and how to handle upstream inconsistencies using pattern-based transformations.

Within a few weeks, the team aligned on a **canonical schema**, reduced duplicate tables by **80%**, and migrated key dashboards to the standardized curated layer—improving both accuracy and refresh performance.

The customer's Analytics Director told me that I “**filled a critical data gap and raised the technical bar for the team,**” and it helped them adopt a **scalable, predictable, engineering-led approach to data.**

Simplifying for customer - Pipeline self-service

“We made the pipeline self-service by moving the logic out of ad-hoc code and into simple, controlled configuration that other teams could own.

In my last project, the data team was getting constant requests to tweak filters, joins, and aggregations. Every change meant a Jira ticket and another code deploy. To fix that, I:

Separated logic from code – we moved business rules into YAML configs instead of hard-coding them in Spark jobs. Things like which columns to include, filter conditions, and basic aggregations were all config-driven.

Standardized patterns – created a small set of reusable templates for common pipelines, so new datasets could plug into the same framework instead of starting from scratch.

Added guardrails – we validated configs, logged changes, and monitored runs so users could experiment safely without breaking production.

Documented in one place – a simple catalog page showing each pipeline, its config, owner, and last run status.

The result was that analysts and data scientists could adjust how their data was processed without waiting on engineering every time, and the core platform team just maintained the framework and guardrails instead of hand-building every pipeline.”**

Machine Learning confidence?

I have not built machine learning models,, but I'm very comfortable with the ML lifecycle because most of my work has been on the data and feature-engineering side. I've partnered closely with Data Science teams to build training datasets, create reusable features, and make pipelines reproducible for both training and serving.

I'm strong at enabling ML teams and designing the data and feature layer that makes their models work in production.

Comfort with databricks stack?

I'm very comfortable with the Databricks stack, both conceptually and hands-on.

In my recent Databricks certification work, I spent a lot of time building on the platform—working with Unity Catalog, understanding its governance model, implementing lakehouse architecture patterns, and working through Delta tables, ingestion, and basic LakeFlow concepts.

Beyond the certification, **I've implemented lakehouse-style architectures for customers using Spark, Iceberg, and cloud-native services, so the design principles—open formats, separation of storage and compute, versioning, governance—are very familiar.** The Databricks way of operationalizing that aligns naturally with how I already build platforms.

My deeper expertise today is in AWS data solutions—Glue, Athena, EMR, Redshift, Kinesis—but over the **past few months I've been intentionally leveling up on the broader Databricks stack**, including security models, lineage, performance tuning, and streaming patterns.

Overall, I'm confident in the platform, I understand how to apply the architecture in real customer scenarios, and I'm excited to go deeper as a Solutions Architect and guide customers on best practices across the Databricks ecosystem."

Top Strength

My top strength is **being deeply customer-focused and building trust with even the most demanding stakeholders.** I've learned that strong architecture only works when people trust you enough to adopt it, so a big part of my strength is connecting with customers, understanding what matters to them, and showing value consistently.

One example is a very senior stakeholder at a large technology company—someone who had been there for more than **30 years** and was known for being extremely tough to satisfy. When I

took over the engagement, there was already frustration on both sides. I made it a point to **establish predictable communication**, set clear expectations, and demonstrate progress every week through architecture walkthroughs, data previews, and short value cycles.

Over time, that consistency changed the relationship completely. The customer began to see me as a trusted advisor, and eventually recognized me with their **Customer Impact Award** for advancing data engineering practices across their team. For me, that validated my belief: trust is built in small moments, and it's one of the biggest drivers of successful delivery.

Weakness?

Balancing Technical Depth with Executive Communication

I come from a strong technical background, so in early stakeholder discussions I used to go too deep into technical detail.

I've learned to adjust my communication style based on the audience — stay high-level with execs and go deep only when needed.

It's something I keep refining, and now I often help junior architects learn how to simplify their message for leadership, which has turned this into a growth area for me.

Future: see yourself in 5 years from now?

In five years, I see myself as a **trusted technical advisor to our customers** — who drive **architecture and strategy for multiple customers** and help them adapt to data and AI solutions of databricks for their business outcomes.

I want to go much deeper in the Databricks stack — Lakeflow declarative architecture, data governance, and the emerging Data Intelligence capabilities — so that **when customers have complex problems or ambiguous requirements, I'm the person they reach out to first.**

I'm also excited about contributing to **reusable accelerators and best-practice patterns** that make our customers more successful and help other architects internally.

Long-term, I'd love to grow into someone who shapes strategy across multiple accounts, mentors newer SAs, and plays a meaningful role in how customers adopt data + AI on Databricks.

What does day to day look like?

“My day-to-day **isn't the same every day**, but most weeks follow a consistent rhythm.

A big part of my role is spending time with client stakeholders—understanding their pain points, clarifying requirements, and helping them think through architecture, governance, and data-quality challenges. I usually take those conversations back and translate them into concrete designs or small proof-of-concepts, especially around Spark, Iceberg, and lakehouse patterns.

I also run working sessions where I present architecture options, walk through trade-offs on cost and performance, and help the customer align in a direction.

In parallel, I work closely with **our internal Slalom team—either shaping RFP responses**, preparing client workshops, or reviewing solution approaches for in-flight projects. And day-to-day I'm partnering with engineers to unblock implementation details, validate pipelines, and make sure what we're building actually ties back to the business need.

Gen AI Exp

I've been working directly with GenAI on AWS through an internal GenAI Bootcamp. As part of that, my team built a capstone project (a multi-state DMV exam study assistant) using:

- LLMs via AWS Bedrock
- RAG architecture with vector databases / Bedrock Knowledge Bases
- Data prep pipelines for GenAI use cases
- Python, LangChain, boto3, Jupyter, and Streamlit

The focus was very much on end-to-end, real-world GenAI systems rather than demos — from data ingestion to retrieval, prompting, and UI.

This kind of pod-based delivery model and backlog of GenAI use cases does sound aligned with my background. Happy to chat more if you'd like to go deeper on specifics or expectations for the role.

DBx Hands on Exp (Accelerator)

I worked on Slalom's **Insurance 360 accelerator**, which is being built on a **Databricks Lakehouse architecture** to unify policy, billing, claims, and engagement data into a governed Customer 360. The design uses a **medallion pattern**, with curated gold-layer customer and policy entities intended to be reusable across analytics, AI, and activation use cases. One of the core capabilities the accelerator is designed to support is **proactive retention** — identifying early churn signals such as missed payments, declining engagement, coverage changes, premium increases, and claims friction, and surfacing them as explainable risk scores in the Customer 360. The intent is to enable insurers to trigger timely, targeted outreach through agents or digital channels, reducing policy lapse risk while improving customer experience and operational focus.

=====

HM round Tips:

1. They will spend 5 mins in intro
2. Behavioural questions - 15 mins
3. Architecture tradeoff
4. Customer presence
5. Machine learning - basics or not - 5 min
6. Why databricks

Tips for intro:

- Show customer readiness
- Executive level communication
- Make it tight
- 45 sec rule - max 1 min
 - Who am I - 10 sec
 - What do I do currently ? **Not history - align with job description**
 - Work directly with whom?
 - Talk about your pre-sales exp
 - HM - Why databricks? - 10 sec

Answering questions -

- Behavioural - Do not go more than 2 mins. Best (1-2 mins)
- For every answer you should have impact - for example this improved x metric, reduce y metrics.
- Leadership (2-3 mins), architecture (2 mins), machine learning questions

- Take time - say that - **I may take a couple of seconds to gather my thoughts, is it okay?** (when gathering thoughts, speak them all... i may speak aloud)
- Do not be very story style type
- When a customer disagrees, you should show empathy, not defend yourself. You should say - I started with fully understanding their concerns. Most of the time it is cost. , complexity, timeline, governance.
- Mention your business constraint - i compared the cost impact of governance with long term flexibility and then some data that you brought...with bench marks, presented them the cost model. This became objective, and not opinionated.
- My goal was not to win the agreement but to help the customer win.
-
- Format:
 - Customer pain point
 - My recommendation
 - Measurable impact
-
- 90 secs - max 2 min rule
- (adding data, keeping time limit)

Follow up questions:

What cloud native services you have used?

What kind of product

How you interacted with business people - i did training, pocs, eng people training.

Behavioural -

Metrics

Deduplication percentage

20% cost reduced - operation overhead.

Aligned with the six architecture teams

Accelerated the decision making by 4 weeks. Spoke with 3-4 business people.

- Have at least 3-4 numbers.

Question - how much confident are you on machine learning thing?

I have not build machine learning models, but I am very comfortable with ML lifecycle, as I have worked with teams at customer side and did feature engineering for them.

Q: your approach to learning new things

- a. I am mostly self-driven. I learn by myself. I follow certain people and (name few people)

-

Top strength and weakness - at 1:09 hour

Strength - more on business side. Scalable , measurable architecture

I influence customer through my trade off thinking.

Maybe learn and grow, always

Weakness:

“A weakness I’ve been actively working on is going too deep into architectural detail too early in conversations.

Early in my career, I would jump straight into patterns, layers, and trade-offs because I was excited about the problem. But I learned that not every audience needs that level of depth.

Over the last year, I’ve been intentional about stepping back first — understanding who’s in the room, starting with outcomes, and only going deeper when it’s relevant. With executives I stay at the value and decision layer, and with engineering teams I dive into the technical path.

It’s helped me communicate more effectively and keep conversations aligned with business goals while still bringing depth when needed.”

Detailed Experience

Final EXPerience / Stories

Intro (45 sec - 1 min mins)

Salesforce Exp (GIC Analytics) - Priority 1:

About the Client / Project

Challenges & Needs

My Role: Solutions Architect / Data Strategist

1. Extending the Engagement through built in trust for being the SME for Data Strategy and Architecture
2. Considered as a trusted advisor for North Star Architecture Alignment
3. Considered as a trusted Advisor for Data Governance
4. Coaching the Salesforce GIC Analytics Team and Bridging Capability Gaps
5. Acting with Urgency to Fill a Skill Gap and Maintain Customer Confidence

Measure/ Quantify:

Google Project

FruitCo Exp (AI/ML Ops) - (Priority 2):

About the client/project

AI/ML Ops:

My Role: Solutions Architect

6. Turning Around an At-Risk Engagement
7. Using Iceberg to Fix Data Quality and Build Trust
8. Account Expansion and Extension at FruitCo
9. Scaling the Architecture to Upstream Data Engineering / DataOps

Measure/Quantify

Apple Exp (Sales Ops):

About the client/project

My Role: Solutions Architect

Measure/Quantify

Results(Impact):

RFPs

BayRen - Priority 4

About the Client / Project

10. Contributed directly to \$2–3M in new D&A pipeline through RFPS and account extensions.

My RFP Role: Technical solutionist

11. Partner Collaboration & AWS MAP Funding (Pre-Sales, AWS)

RFP - Salesforce Informatica Merger - Priority 4

Key Challenges

Irhythm - Workshop & POC (Pre-Sales) - Priority 3

About iRhythm

Challenges

My Role: Solutions Architect

Databricks - Slalom Partnership

13. Databricks Brickbuilder Program Participation (Databricks - Pre sales)

14. Databricks - COE - accelerators and certifications

Open AI - Technology partners -partnership (pre-sales)

15. Generative AI Pilot (OpenAI Codex):

CMD5 - Multiple Technology Partners - Snowflake, Fivetran, Dataiku, ThoughtSpot, and Alation.

16. Slalom CMD5 Partnership with Snowflake and Ecosystem Tools (Multiple clients - Pre sales)

Learning and being curious

18. Constantly Learning and Staying Ahead (Databricks & RAG Architecture)

ICE Mortgage

About the client/project & challenges

My Role: Data Architect / Data Engineer

19. Performance Optimization with EMR (ICE Product)

LPs mapping & Keywords

Intro (45 sec - 1 min mins)

I'm a **Solutions Architect in the Data & Analytics practice, working with some of our Slalom's strategic NorCal accounts.** I have been with Slalom for around 3 years now, and based out of the East bay. My focus has been **advising, guiding, and building scalable data solutions** for customers — my work spans across **both pre-sales and delivery tracks.**

On the pre-sales side, I support

- **RFPs by providing technical solutions,**
- **run workshops** with new and existing clients to **generate leads and opportunities,** and
- **collaborate with our technology alliance managers to promote partner platforms, grow the customer base, and create new opportunities.**

On the delivery side, I've recently led **several architecture engagements.**

At Salesforce, I

- **led the architecture for their Global Incentive Compensation analytics team** — making sure it aligned with their North Star vision and enterprise standards. In the process,
- **I partnered with several business and technical teams to define how the data from different systems will flow into Snowflake** in a way that was both consumable for analytics and compliant with SOX.

For FruitCo— a large consumer electronics and technology company —

- I led the design and architecture of their AI/ML operations lakehouse on AWS and Iceberg, supporting large-scale model-training and data-quality workflows.

Before Slalom,

- I led the data platform team at ICE Mortgage Technology in Pleasanton, where I worked with **production workloads** — **optimizing EMR jobs, optimizing spark performance,** and modernizing pipelines.

I am also a trusted AWS architecture advisor in our D&A practice —

- **helping with sales pursuits and building our team's skills in the D&A space.** Lately, I've been part of Databricks learning cohorts and GTM enablement sessions with our Databricks partners.

So that's a bit about me and the kind of data architecture and solutioning work I've been leading across clients. Happy to share more details as needed.

I lead the design and delivery of end-to-end data platforms for enterprise clients.

Salesforce Exp (GIC Analytics) - Priority 1:

About the Client / Project

- **One of my customers** is a large global CRM platform provider. I work with their **Global Incentive Compensation (GIC)** team, who are
 - responsible for building and maintaining the **commission logic—the rules that determine how sellers get paid.**
 - Their main challenge was ensuring that every seller's **compensation statement** was accurate, despite having very complex and frequently changing incentive structures.

Challenges & Needs

- The compensation data is gathered from multiple systems like Workday, and other native apps
- The Analytics team relied on the **BT group for data**, but since **BT had global priorities, their timelines often didn't align with the GIC team's needs.**
- They were facing failed audits from the **SOX compliance team** and were spending a lot of time on manual audit tasks—like pulling reports, tracking access, and validating controls.
- Overall, leadership needed **reliable, automated, and SOX-compliant reporting to gain trust and accuracy in the compensation process.**

Slalom's Role:

- Slalom was engaged to perform a **discovery phase, assess the current state, and define future state recommendations.**

My Role: Solutions Architect / Data Strategist

Result:

1. Extending the Engagement through built in trust for being the SME for Data Strategy and Architecture

Actions:

- I led a Slalom team of 2 members in the initial phase to do the data discovery and current state assessment which required:
 - **interviewing stakeholders across teams** and understanding their needs and pain points.
 - **Reviewing existing data assets and infrastructure setup and pipeline**
 - **talking to both the analytics and SOX teams** to understand pain points and review earlier audit reports.
- I mapped out the data issues, **showed the real business impact (33% of the data is only provided by DET in SOX compliant way)**, and helped leadership see where the biggest gaps were — which got everyone aligned on improving both process and architecture.
- I recommended future-state data architecture options — **S3 + Snowflake and an S3-based lakehouse** — and **walked the team through the trade-offs of each.**

Result2:

2. Considered as a trusted advisor for North Star Architecture Alignment

Actions:

- I looked into Salesforce's global technology team (DET) technical architecture and **met with enterprise architects to understand the overall architecture, and design an architecture that fit into the company's overall direction**, not just the GIC analytics setup.
 - The goal was to **build on the existing Snowflake instances and create a solution that could scale gradually without disrupting business operations**, while paving the way for a more stable and improved data pipeline.
 - Ownership of data
 - Defining data flows their key data assets

- **I worked with executives to align on the GIC Analytics North Star vision and built an architecture that aligned with the vision and made compensation data more reliable, accurate, and timely.**
 - **For stable, long-lived SOX reports, built and maintained by DET in the certified EDW**
 - **Utilizes a "Controlled DEP" for evolving SOX controls, ad-hoc audit requests, strategic decision support and compensation modeling**
- **I presented the architecture to the SOX and DET teams to ensure alignment on how it would integrate with upstream data systems.**

Result 3:

3. Considered as a trusted Advisor for Data Governance

Actions:

- Over time, the customer trusted me more and started depending on me to shape the data governance approach and make sure it fit into the overall architecture and roadmap.
- **Reviewed the existing access controls** and defined the improvements **using the RBAC in Snowflake**, as well as defining data ownership
- Removing the ambiguity of ownership, and stewardship of data

Result 4:

4. Coaching the Salesforce GIC Analytics Team and Bridging Capability Gaps

Actions:

- created **detailed technical specs explaining each layer** of the architecture — why it existed, how data should flow through it, and what best practices to follow for scalability and governance.
- held **several working sessions and walkthroughs** with their data engineers to explain implementation details, answer questions, and reinforce best architecture practices
- The client's analytics director had mentioned that I **helped fill a critical data gap** and raised the technical bar for the team.

Result 5:

5. Acting with Urgency to Fill a Skill Gap and Maintain Customer Confidence

Actions:

- Mid phase, **The customer's expectations had shifted — the role had become much more project-management heavy**, requiring strong coordination and proactive communication.
- I worked **closely with our account and delivery leadership** to assess the situation and quickly align on the right profile for replacement.
- I moved fast to hire someone who was a better fit for the role's new scope, while making sure the transition didn't impact delivery.

Measure/ Quantify:

- Data from **10 different sources systems like Salesforce, Workday, and other native apps.**
- 20+ different compensation programs
- Documented the **80+ reports and their sources**
- Co-ordinated/ collaborated with **2 BT teams, 3 Upstream Teams, 2 Enterprise architects, and Managers and SOX compliance team**
- Coordination with Account team people and Slalom leadership
- Presented to 5 different groups

Google Project

- On the Google WPA project, I established the technical foundation by designing the end-to-end data architecture, resolving PII and security requirements for the Google WEPS UX team, and delivering a clear requirements document and solution design that aligned all stakeholders on the implementation path. I operated with minimal direction, translated ambiguity into structured deliverables, and built strong rapport with Hong and Sarah. My technical solutioning, productivity, and ability to communicate across technical and business audiences were key in moving the project forward despite infrastructure and dependency challenges.

FruitCo Exp (AI/ML Ops) - (Priority 2):

About the client/project

AI/ML Ops:

FruitCo's AI & Machine Learning (AIML) organization **powers intelligent features** across core applications such as Siri, Mail, Calendar, Music, and Photos.

Within this group, the **Annotation Operations (Annotation Ops) team** provides the human-in-the-loop foundation for Fruitco's machine learning — a 6,000-person global workforce responsible for labeling and validating data that trains Apple's ML models.

Data-Driven Operations: Productivity, quality, and quota compliance are tracked continuously through dashboards and **performance analytics**.

Challenges:

- The **existing data ingestion process was not scalable or standardized**, leading to frequent failures and manual interventions.
- Multiple disjointed pipelines and embedded transformation logic within Tableau made reporting unreliable and hard to maintain.
- The upstream data provider teams were not aligned or supportive in delivering data in the required structure and cadence for the BI team.

My Role: Solutions Architect

6. Turning Around an At-Risk Engagement

- Brought into an at-risk engagement after consultant left mid-project
- **Discovery phase was complete but** lacked direction and next steps
- Partnered with orgs **DataOps team** to clarify requirements and infrastructure needs
- Identified key data sources to ingest and curate first
- Proposed **fast-track MVP: start with one dataset to show quick value and set foundation for scale**
- Delivered MVP successfully — restored client confidence and rebuilt trust in the account

7. Using Iceberg to Fix Data Quality and Build Trust

- Designed **data-engineering approach using Apache Iceberg for time travel, incremental refresh, and upstream update tracking**
- Implemented **schema-evolution standards** to handle upstream data changes smoothly
- **Improved data quality, traceability, and reliability** through Iceberg's metadata and versioning

8. Account Expansion and Extension at FruitCo

- Built a **self-service analytics pipeline** enabling business users to explore and onboard datasets independently
- Focused on scalable ingestion and transformation pipelines with **downstream analytics and ML use in mind for around 6000 or so global workforce doing 100 annotations per day**.

- **Collaboration with data scientists** to curate datasets at the right granularity and privacy standards
- Outcome: **engagement expanded into a larger, long-term partnership** for continued platform scaling
- Received a **Customer Impact Award from a senior, long-time Apple stakeholder in recognition of the results** (you can start with - i do want to mention)

9. Scaling the Architecture to Upstream Data Engineering / DataOps

- Integrated data workflows with MLOps pipelines for automated retraining and deployment via Airflow and CI/CD
- Designed standards for incremental data refresh and tracking of upstream changes adopted by DataOps team
- Established foundation for scalable, governed data platform supporting both analytics and ML initiatives
-

Measure/Quantify

- 6,000-person global workforce
- 100 annotations daily by an analyst
-

Apple Exp (Sales Ops):

About the client/project

Challenges:

My Role: Solutions Architect

- Apple sales - I led the data-harmonization workstream — met with teams across retail, digital, and carrier sales, documented pain points, and defined how data would be ingested and transformed.

Measure/Quantify

-

Results(Impact):

RFPs

BayRen - Priority 4

About the Client / Project

- BayREN manages six separate energy and water efficiency programs across the Bay Area.
- Each program had its own database — **no central repository or unified reporting.**
- The goal was to build a **Data Aggregation and Reporting Tool** to automatically pull, transform, and consolidate data across all six programs into a single reporting layer.

Key Challenges

1. **Six disconnected databases with different schemas** and data formats (**PostgreSQL, Excel, and CSV exports**).
2. Manual, time-intensive reporting with inconsistent metrics and no automation.
3. Needed a common data model and reliable refresh mechanism to generate unified metrics and dashboards.

10. Contributed directly to \$2–3M in new D&A pipeline through RFPS and account extensions.

My RFP Role: Technical solutionist

- Worked with **client executives, practice leads, and the pursuit team** to define requirements and shape the technical solution.
- Wrote the **technical sections of the RFP**, including:
 - Approach for the Technical Solution
 - Example of a Cloud-Hosted Data Platform on AWS
 - Assumptions and Constraints in Selecting the Solution
- Designed a **scalable AWS data architecture** using **S3, Glue, Redshift, and Lake Formation**, ensuring it met both governance and scalability needs.

- Helped draft the **SOW, RACI**, and **delivery plan**, aligning technical scope, staffing, and pricing with the proposal.
- Joined the **orals presentation** to walk the both technical and business audiences on the customer side through the proposed architecture and delivery approach.
- Contributed to **\$2–3M in new pipeline** for the Data & Analytics practice.

11. Partner Collaboration & AWS MAP Funding (Pre-Sales, AWS)

“We had a client who wanted to test a new data architecture, but they didn’t have any budget for a **proof of concept**, and we didn’t want to lose momentum.

I was asked to find a way to fund the POC so we could move forward and still show quick value.

I reached out to our **AWS alliance team** and explored options through the **MAP program** — the Migration Acceleration Program.

I drafted a short proposal explaining the use case, the AWS services involved, and how it could increase their cloud consumption over time.

Once the funding was approved, I led a small **cross-practice team** — data engineers — to run the POC and integrate **with AWS analytics services**.

The POC was completely funded by AWS, so the client didn’t have to pay a thing but still got a working demo in their environment.

That success built trust with the client, opened the door for a larger engagement, and also strengthened our partnership with AWS for future projects.”

RFP - Salesforce Informatica Merger - Priority 4

About:

- Objective: **plan the integration of Informatica systems and data into Salesforce’s ecosystem post-merger**
- Focus on system mapping, data flow consolidation, and risk-aware migration roadmap

Key Challenges

- **Large number of overlapping systems and data integrations** between both companies
- Lack of unified inventory of critical systems and dependencies
- High **risk of business disruption** during migration without clear prioritization

- Need for a **multi-phase roadmap** balancing quick wins and long-term stability

My Role:

- Proposed **migration-readiness assessment** using structured scoring (complexity, data volume, business criticality)
- Documented **risks, assumptions, and mitigation plans** for inclusion in the RFP response
- Approach for below outcomes:
 - **Identify and prioritize the most critical systems, data flows, and integrations for migration to support core business processes**
 - Reduce risk by flagging **migration challenges and high-impact dependencies** early
 - Accelerate implementation with a **migration readiness report** and multi-horizon migration roadmap

Salesforce is seeking proposals from qualified consulting firms to support Informatica M&A integration planning program.

Systems and Data Flow Assessment - detailed out assessment approach and acceleration methods

Irhythm - Workshop & POC (Pre-Sales) - Priority 3

About iRhythm

- iRhythm is a health-tech company developing **wearable cardiac-monitoring devices** to detect heart rhythm disorders.
- This was a **new logo** for Slalom, so it was important for us to make a first impression.
- Their **Data Systems & Services group** needed to generate faster, more reliable insights from the **vast amount of device data** being collected — to improve both operational efficiency and patient outcomes.

Challenges

- At the time, iRhythm was in the middle of a major platform transition:
 - They were **moving their data warehouse from Redshift to Snowflake**.
 - Data ingestion relied on new **Go-based extractors**, still in development.
 - There was **no established ETL/ELT framework**, which made data movement, transformation, and validation slow and inconsistent.

My Role: Solutions Architect

Led a **two-week discovery** and assessment of iRhythm's data platform.

12. Data Discovery & Demo Workshop Leading to Extended Engagement (iRhythm - Pre-Sales)

- Proposed and delivered a **four-week test-and-learn POC** to demonstrate:
 - Automated **data discovery** and cataloging.
 - **PII-aware metadata tagging** and privacy compliance.
 - **Data-quality monitoring** across representative datasets.
- Designed and demoed the solution using **AWS Glue** (cataloging, schema tracking) and **Amazon Macie** (PII detection), integrated with quality rules and remediation workflows.
- Presented the architecture to both **executive and engineering teams**, resulting in the acceptance and leading it to extended engagement
-

Measure/ Quantify:

- Tea size: 6
- Duration: 6 weeks (2 week discovery, 4 weeks, test and learn)

Databricks - Slalom Partnership

- With growing client demand for modern data platforms, Slalom has been deepening its co-selling and solution-development partnership with Databricks. The focus is on building industry accelerators and repeatable Lakehouse frameworks that quickly demonstrate business value.

My Role

- Partner Enablement & Accelerator Development
- I've been collaborating closely with Databricks partner managers to identify high-impact industry accelerators — focused on Lakehouse and Unity Catalog — and to enhance existing assets or design new ones aligned to customer use cases.
- These accelerators aim to shorten delivery time, improve governance, and help clients adopt the Databricks Lakehouse more effectively.
-

BrickBuilder Series

-
- I participated in the Databricks BrickBuilder program, which strengthens go-to-market alignment and co-selling opportunities between Databricks and its partners.
- The program showcases partner-developed solution assets — either horizontal technical solutions or industry-specific offerings — and includes partner-led enablement sessions where Databricks architects share best practices for Delta Lake, Unity Catalog, and Lakehouse governance.
- I also earned Databricks badges through this program, deepening my technical proficiency in these areas.
-

Data Engineering Cohort Development

- Participating in Slalom's internal Data Engineering cohort to upskill teams on Databricks Lakehouse architecture, earning Data Engineer and pursuing professional certifications to enable future Databricks-based client engagements.

13. Databricks Brickbuilder Program Participation (Databricks - Pre sales)

14. Databricks - COE - accelerators and certifications

Open AI - Technology partners -partnership (pre-sales)

About - Slalom is participating in an **early pilot with OpenAI's Codex**, an AI-powered code composition and agentic development tool. A small internal cohort (including you) will be granted Codex licenses to **experiment, provide structured feedback, and shape Slalom's long-term adoption strategy** for generative AI in software and data engineering.

15. Generative AI Pilot (OpenAI Codex):

- Participating in an early internal pilot with OpenAI's Codex, an AI-powered code composition and agentic development tool. As part of a small cohort, I'm experimenting with its applications in data engineering and solution development, providing structured feedback to shape Slalom's generative AI adoption strategy.

CMDS - Multiple Technology Partners - Snowflake, Fivetran, Dataiku, ThoughtSpot, and Alation.

16. Slalom CMDS Partnership with Snowflake and Ecosystem Tools (Multiple clients - Pre sales)

“As part of Slalom’s **Cloud Modern Data Stack (CMDS)** initiative, I worked on a partnership effort with **Snowflake** and several ecosystem vendors — including **Fivetran, Dataiku, ThoughtSpot, and Alation**.

The goal was to design and present a **modern data-architecture blueprint** that showed how these partner tools could work together to accelerate data modernization for our clients.

I collaborated with our Snowflake partner team and vendor solution architects to evaluate how each tool fit into the end-to-end data flow — ingestion, transformation, governance, and visualization.

We defined reference architectures that showcased **Fivetran for ingestion, Snowflake for storage and compute, Dataiku for machine learning and data science workflows, Alation for cataloging and governance, and ThoughtSpot for self-service analytics**.

I helped document and present this architecture to both internal teams and prospective clients, positioning it as a **ready-to-deploy accelerator** for cloud data modernization.

It not only demonstrated how these technologies complement each other but also **helped Slalom and Snowflake co-sell** the joint value proposition to clients.”

17. OpenAI partnership - codex tool evaluation

Learning and being curious

18. Constantly Learning and Staying Ahead (Databricks & RAG Architecture)

“With more of our clients asking for **Databricks-based solutions**, I’ve been very intentional about **continuing to upskill** and stay ahead of what’s coming next in the data space.

I recently completed several **Databricks certifications** and participated in **partner-led enablement sessions** focused on Lakehouse architecture, Delta Live Tables, and governance with Unity Catalog.

It's helped me better understand how Databricks positions its platform end-to-end — not just for analytics, but also for machine learning and real-time data engineering.

At the same time, I joined an internal **learning cohort** where we explored **Retrieval-Augmented Generation (RAG)** architecture — combining vector search, LLMs, and Databricks' AI capabilities.

I worked hands-on to build a small **proof of concept** using this pattern and later **presented it to our internal team** to share learnings and discuss how we could extend it for client use cases.

For me, continuous learning is part of the job — it's what helps me bring fresh, relevant perspectives to client conversations and stay aligned with how the data and AI landscape is evolving.”

1.

ICE Mortgage

About the client/project & challenges

ICE Mortgage's **Data Connect** transforms complex, high-volume Encompass loan data into **structured, queryable datasets** for lenders — but faces challenges around **real-time data consistency, schema evolution, massive scale, and multi-tenant orchestration** while maintaining performance and reliability.

My Role: Data Architect / Data Engineer

Stories:

19. Performance Optimization with EMR (ICE Product)

“While working on the **ICE data product** for one of our major clients, we started noticing performance bottlenecks in their reporting layer — queries were running slowly, and the customer was getting frustrated with long refresh times.

I dug into the issue and realized that some of the large-scale transformations were running inefficiently on their existing setup.

To address this, I proposed moving the heavy workloads to **AWS EMR**, where we could take advantage of distributed processing and better resource tuning.

I worked with the team to reconfigure the data pipelines, optimize Spark jobs, and fine-tune the cluster settings to balance performance and cost.

Once deployed, the **query times improved by nearly 40%**, and the customer immediately noticed the difference.

This optimization not only made the ICE product much faster and more reliable, but it also **earned strong trust from the customer**, who started involving us in deeper discussions about scaling their platform and performance strategy going forward.”

- a. Drove **60% cost reduction** in EMR streaming workloads and **80% faster Athena query performance** through Spark tuning and architectural redesign.
 - b. Implemented **observability frameworks** (Spark UI in private subnet, EMR monitoring) and benchmarked all critical releases to ensure SLA compliance.
20. Architected the Enterprise Data Platform at Ellie Mae (ICE)
- a. Created a **generic ingestion architecture** migrating from RabbitMQ to Kafka, enabling scalable, real-time data delivery.
 - o Acted as bridge between engineering, QA, and product management teams to align technical design with business objectives.

LPs mapping & Keywords

📌 Keywords-LPs-Questions

Team Moghul Award

You were nominated for Team Mogul NorCal H2. Below nomination information!

From SD Singh "Salesforce ICM Modernization Team":

"Slalom team is an integral part of the SPIFF Implementation team at Salesforce. Has been able to address key capability gaps to achieve program goals including experience design, process re-engineering, SOX program management as well as workstream and program leadership. Slalom team members have demonstrated excellence in delivery and subject matter knowledge to assist Salesforce leaders in implementing key capabilities related to compensation planning and operations, next-generation analytics and payee experience."

ARCHIVED

ARCHIVED

Other Behavioural Questions

Other Behavioural Questions

1. Tell me about yourself	1
2. What is your top strength?	2
3. What is your weakness?	2
Intro (2 mins)	3
✓ 2. How does your day-to-day look like?	4
✓ 3. Salary expectations (the safest framing)	5
Questions	5
Mind Map	5

1. Tell me about yourself

◀ Past:

I started my career as a data engineer, working hands-on with pipelines, data modeling, and analytics delivery. Over time, I moved into architecture because I found myself naturally connecting business goals with technical solutions — understanding not just how to build something, but why it matters to the end users.

▮ Present:

Today, as a Senior Solutions Architect in Slalom's Data & Analytics practice, I work across both pre-sales and delivery. I create technical solution documents, lead client discovery workshops, and coordinate between account teams, engineers, and partner architects — especially for AWS-based data platforms.

Recently I've been involved with a large CRM company to design an incentive analytics architecture, and a consumer-electronics client to unify their global sales data. My day-to-day revolves around aligning stakeholders, documenting current-state gaps, and defining architectures that are feasible and governed from the start.

▶ Future:

Going forward, I want to continue in architecture roles that bridge strategy and implementation — working on complex, multi-team data ecosystems. I'm also investing time in deepening my Databricks and AI expertise.

Keywords: data architecture, AWS, cross-functional coordination, discovery, RFPs, governance, practical solutions.

2. How is your day going so far?

Ans: It's going great, thanks for asking! I've been really looking forward to our chat today - how's your day going so far?

3. Where do you see yourself in 5 years?

Intro (2 mins)

I'm a **Solutions Architect in the Data & Analytics practice, working with some of our Slalom's strategic NorCal accounts**. I have been with Slalom for around 3 years now, and based out of the East bay. My focus has been **advising, guiding, and building scalable data solutions** for customers — my work spans across **both pre-sales and delivery tracks**.

On the pre-sales side, I support

- **RFPs by providing technical solutions**,
- run **workshops** with new and existing clients to **generate leads and opportunities**, and
- **collaborate with our technology alliance managers to promote partner platforms, grow the customer base, and create new opportunities**.

On the delivery side, I've recently led **several architecture engagements**.

At Salesforce, I

- **led the architecture for their Global Incentive Compensation analytics team** — making sure it aligned with their North Star vision and enterprise standards. In the process,
- **I partnered with several business and technical teams to define how the data from different systems will flow into Snowflake** in a way that was both consumable for analytics and compliant with SOX.

For **FruitCo**— a large consumer electronics and technology company —

- **I led the design and architecture of their AI/ML operations lakehouse on AWS and Iceberg**, supporting large-scale model-training and data-quality workflows.

Before Slalom,

- I led the data platform team at ICE Mortgage Technology in Pleasanton, where I worked with **production workloads** — **optimizing EMR jobs, optimizing spark performance**, and modernizing pipelines.

I am also a trusted AWS architecture advisor in our D&A practice —

- **helping with sales pursuits and building our team's skills in the D&A space.** Lately, I've been part of Databricks learning cohorts and GTM enablement sessions with our Databricks partners.

So that's a bit about me and the kind of data architecture and solutioning work I've been leading across clients. Happy to share more details as needed.

-

2. How does your day-to-day look like?

Version to Say Out Loud (30–40 sec)

“My days are a mix of pre-sales architecture work, hands-on technical solutioning, and partner collaboration. I support Slalom pursuit teams on in-flight deals — doing discovery, designing architectures, running workshops, or building demos.

With dedicated customers, I work closely with their engineering and data teams, helping them evaluate current-state platforms and shaping the right modern data solution — often across AWS, Databricks, or Snowflake.

I also spend a meaningful amount of time learning and upskilling — I'm part of the Databricks cohort at Slalom and actively working on Lakehouse certifications. So overall, it's a balance of customer-facing work, technical deep dives, and continuous learning.”

-

3. Salary expectations (the safest framing)

“I’m very flexible. For me, the most important thing is the right role, team, and growth path. I trust that Databricks offers competitive, top-of-market compensation for its Solutions Architect roles, so I’m comfortable aligning with your bands once we determine that it’s a mutual fit.”

Use this to maintain boundaries while staying cooperative:

“I completely understand. My priority is finding the right role and team. I’m confident we can land on something fair once we move further in the process.”

Questions

1. “What does success look like for a Solutions Architect on the Digital Native team in the first 90 days?”

“Before we wrap, I’d love to get your guidance on what to expect in the upcoming interview rounds so I can prepare effectively. What does the interview sequence typically look like for this team?”

=====

Mind Map

1. Give me an example of a time you used customer feedback to drive improvement or innovation. What was the situation and what action did you take?
2. Tell me about the most innovative thing you’ve done and why you thought it was innovative (can also probe with: That sounds more evolutionary than revolutionary – tell me about something you’ve done you feel was truly revolutionary? Ask for one or two additional examples to see if it’s a one off or pattern.)
3. **Tell me about a time you were able to make something significantly simpler for customers. What drove you to implement this change?**
Talk about self service pipeline

Apply:

Newer:

You will guide customers in designing advanced cloud-based data architectures, implementing innovative technologies across AWS services (including Amazon RDS, Aurora, Redshift, Athena, EMR, Glue, Kinesis, MSK), partner solutions (Snowflake, Databricks), and open-source frameworks (Spark, Flink, Trino), while adopting the best cloud practices for databases, storage, and analytics services.

[Sr Data Solutions Architect, Strategic Accounts - AWS](#)

https://www.databricks.com/company/careers/field-engineering/solutions-architect-retail-6456675002?gh_jid=6456675002&gh_src=62a881d62

https://www.databricks.com/company/careers/-field-engineering---indirect/sr-specialist-solutions-architect---aws-partner-solution-architect-8108693002?gh_jid=8108693002&gh_src=62a881d62

https://www.databricks.com/company/careers/field-engineering/solutions-architect-retail-6456675002?gh_jid=6456675002&gh_src=62a881d62

https://www.databricks.com/company/careers/-field-engineering---indirect/sr-specialist-solutions-architect---aws-partner-solution-architect-8108693002?gh_jid=8108693002&gh_src=62a881d62

Older:

Apply:

<https://www.linkedin.com/jobs/collections/recommended/?currentJobId=4321885578&originToLandingJobPostings=4335531262%2C4321885578%2C4317369432%2C4335511113%2C4321739937>

<https://www.linkedin.com/jobs/collections/recommended/?currentJobId=4335511113&originToLandingJobPostings=4335531262%2C4321885578%2C4317369432%2C4335511113%2C4321739937>

<https://www.linkedin.com/jobs/collections/recommended/?currentJobId=4321739937&originToLandingJobPostings=4335531262%2C4321885578%2C4317369432%2C4335511113%2C4321739937>

https://www.linkedin.com/jobs/view/4312177404/?eBP=CwEAAAGac_-qdLS_SSMp94FO58PZoybWDmMecXa2L9Bxn18h9pd7_tY92Xewrs3x6tNn8VqhmGB_hAmn3nuP_HLbej7n2QWXHb1fchhRd196Qfvj9Pu6stx0--CoL-kulQ3feXLfHeOffliE2QfSLowmCW4sSJ84ccMfTqH2T5MJdmJfHno9xiNZ4J9S0o_UV0ZGF5txzS8OVmZMM6BblSF9CfsDAdhiM0PV3hQYj6obWnXpA-GNBVMhwxT_3ASOWn1XrWESJVTW4PcgxvQqA9PcJ7mDIzjD1J8RUyQEzd7sJntphTz9OLXpDjh-5Z-X-3

[xFqTGfagRo8cFYL6ym6ylsqw2HfrUQCQIXmzBHby0N3an7EKGcq_Ry5pyBJfq8rllmeJxIJk55wa7hpgfl56wjnWK-6S3m6L1d3b9QPv270z3p5eccyrKCPkMmPHALzpT11BcxLMxcEyiaTi03efNKd7nHv2ckuhuoOSgzU_x0&trk=flagship3_jobs_discovery_jymbii&refId=gm%2F%2FDjuTS4cdfE2SZ4Cirw%3D%3D&trackingId=8vx7Je755piMOeVb8A08yQ%3D%3D](https://www.linkedin.com/jobs/view/4321885578)

<https://www.linkedin.com/jobs/view/4321885578>

Asheesh:

<https://www.linkedin.com/jobs/collections/recommended/?currentJobId=4321885578&originToLandingJobPostings=4335531262%2C4321885578%2C4317369432%2C4335511113%2C4321739937>

Apply:

https://www.databricks.com/company/careers/field-engineering/solutions-architect---digital-native-business-6117719002?gh_jid=6117719002&gh_src=si050rvy2us

https://www.databricks.com/company/careers/-field-engineering---indirect/sr-specialist-solutions-architect---aws-partner-solution-architect-8108693002?gh_jid=8108693002&gh_src=62a881d62

https://www.databricks.com/company/careers/field-engineering---fe-direct-regulated/solutions-architect-8243221002?gh_jid=8243221002&gh_src=62a881d62

<https://www.linkedin.com/in/ericjonhenderson/>

<https://partner-academy.databricks.com/pages/91/agents-bricks-catalog-partners>

- Expertise in data governance systems and solutions that may span technologies such as Unity Catalog, Alation, Collibra, Purview, etc.
- Experience with high-performance, production data processing systems (batch and streaming) on distributed infrastructure.

Stripe

Stripe

<https://stripe.com/jobs/listing/partner-solutions-architect/7251505>

<https://www.linkedin.com/jobs/view/4318720474>

As a member of the team, the Solutions Architect will build highly consultative trusted relationships with c-level executives and technical decision makers in the customer's organization to provide business and technical thought leadership, and become a trusted advisor in solving mission-critical business challenges. She/he will architect business models with them so they can drive new monetization opportunities and growth for both the user and Stripe.

Responsibilities

- Be a key member of a Stripe account team with deep involvement in defining the customer strategy
- Identify and close new complex opportunities, and accelerate solution adoption. Understand and align with the customer's product and business roadmap and uncover opportunities for Stripe to accelerate their progress on that roadmap.
- Drive thought leadership and be a trusted advisor for key business and technical stakeholders including CIO, CTO, COO, CRO, CDO, Head of Product Engineering and Head of Payments.
- Work to deeply understand the customer's business objectives, and then design and apply Stripe's solutions to their challenges. **Develop and articulate the end to end customer transformation journey and create future roadmaps and architectures with deep 3rd party integrations in play.**
- Conduct business analysis and background research on customers, industries, and competitors with a view to support growth for the user's business.
- Lead strategy engagements with users to align to business priorities, design solutions and secure the technical win by performing in-depth customer discovery, evangelizing the proposed solutions through executive briefings, readouts, presentations, workshops and solution demos.
- Expand Stripe's footprint within existing accounts to new global markets and business units by introducing new capabilities and local payment methods.
- Act as the "Voice of the Customer" to bring insights, trends and opportunities to the leadership, product and engineering teams from the forefront and prioritized for delivery.
- Participate in sales forecasting, Quarterly Business Reviews and account/territory planning to strategically and tactically align for success
- Approximately 25% travel is anticipated to help you connect in-person in a meaningful way with users and collaborate internally

Final Stories-backup

Final Stories

2. What i want to tell : Turned Around an At-Risk Engagement (**Apple**).

S: I was brought in in an urgent distrust situation when a consultant had left. The discovery phase was completed, with no clarity of what work needed to be done

T: I was supposed to quickly understand and stabilize the next phase to design and subsequently lead a team for implementation

A: I did deep to understand the current state, discovery of data that needed to be ingested and curated, and proposed / presented an architecture with all the different key use cases, and their ecosystem, which was well received by the client, along with an approach to move fast and quickly to deliver MVP, with a particular dataset, and leading the team to build an end to end use case.

R: The customer gained more trust in slalom and further, relied on me as SME consulting and looking for solutions for the data issues

Keywords while being interviewed:

LPs:

Questions it can address to:

3. What i want to tell: Iceberg utilization and advanced DE knowledge to handle data issues

S: At the beginning of the last year, the customer was frustrated by unexpected issues and changes stemming from the data sources upstream, which were introducing data quality issues into their leadership reports.

T: Bring in capabilities to address the data quality issues and increase trust in data

A: designed and presented to seek alignment with the team and leaders. Design addressed issues like late arrival of data, time travel, incrementally refresh data and tracking updates of the data upstream . used Iceberg for implementation and

R: The critical datasets were handled specifically for these nuances, and these are now continuously running in production with failure logs. This increased the trust in data for BI analysts and data scientists to use this platform which kept on increasing with more datasets for a longer term. **Customer Impact Award from Apple Executive**

Keywords while being interviewed:

LPs:

Questions it can address to:

4. What I want to tell: Acted as a technical advisor to the data governance workstream of the customer at salesforce.

S: The GIC team was facing a lot of compliance issues w.r.t. SOX compliance, and also they did a lot of manual work providing audit reports to the sox team

T: I was asked by the customer to provide guidance and lead to their data governance needs, and define and suggest best practices for the data access patterns in snowflake, and .. and the roadmap and path for long term, SOX compliance,

A: navigated and discovered the pain points and their existing state of data access and ownership patterns. Spoke to SOX compliance team and reviewed the earlier compliance reports, to understand what is going on. Suggested specific measures like RBAC and granular access controls, with automated audit controls to give trust to sox compliance team, and

R: generate trust in the data pipelines that GIC analytics team was suing. The team was doing less manual efforts to provide manual reports to ensure compliance and increased focus and productivity on data analysis and BI activities of the team

Keywords while being interviewed:

LPs:

Questions it can address to:

5. What I want to tell: Extended the engagement with the customer(**Salesforce**) to help advise on data strategy and meeting long term vision & goal and providing technical solution for the same

S: Slalom was hired for the discovery phase and was finishing up with findings and key gaps and future state recommendations.

T: to provide the best future state recommendation and roadmap based on assessment

A: Based on the discovery we did at customer's eco-system, I went ahead and generated numbers in terms of highlighting the key data gaps, and its impact on business. I led discovery workshops, defined architecture options, and presented data-governance and ingestion frameworks that were later approved by executives. Those conversations helped shape the final engagement

R: The presentation led the customer to be more interested in slalom and an interest to continue the engagement for taking the architectural recommendations further.

Keywords while being interviewed:

LPs:

Questions it can address to:

6. What I want to tell: presented to enterprise level architects and VPs the architecture and roadmap to architecture to meet the north star vision and goals of gic with tradeoff and got alignment.

S: The GIC analytics team was very much dependent on various BT (Business technology) teams to provide the data in shape and form that is useful for them to generate meaningful insights and make critical decisions on sales compensation. There was a delay caused because of misalignment of priorities for BT team and was affecting overall key business decisions and insights in the GIC analytics team. Identified that the team was bringing in their data manually on their own to meet the business needs.

T: i was tasked with sharing the requirements and future state required to meet the north star vision

A: Based on the initial discovery phase, I designed and presented future state scalable architecture to the BT architects emphasizing the need for this architecture and as well VPs to help understand the business value we get out of it.

R: The future state architecture vision helped the VPs make a decision to invest in pieces in the same quarter and plan pieces of this in the coming quarters.

Keywords while being interviewed:

LPs:

Questions it can address to:

7. What I want to tell: Did a data discovery and demo workshop with a health client resulting in larger engagement with the client.

S: iRhythm was a new logo for the slalom team, new client. Slalom was in 3rd week of strategy engagement - lot of moving parts, they don't have full fledged ETL or ELT tool , they were writing their GoLang extractor, they have decided to move from redshift to snowflake. iRhythm's service delivery team has embarked upon a journey toward next-generation data services platform to support the company's rapid growth. Slalom will support the service delivery team to establish strategy and roadmap for the platform and help with execution of one "test & learn" experiment in a non-production environment for the data services platform.

T: i was brought in the test and learn phase, where the primary objective was to do suggest architecture options based on the outcomes from exploration phase, and then a do a poc for end to end for a particular use case in their environment to evaluate and confirm on the architecture.

A: did evaluate architectural options and did poc in their environment, and at the end of 4 week, did a demo presentation to customer on data cataloging service, schema monitoring and privacy catalog using glue in aws. Also used Macie for data quality integration. Purpose:

Central catalog that shows detail about data sources in Snowflake Data Warehouse

This catalog is automatically triggered when data lands in S3 Data Lake

Snowflake uses this data catalog for Data Masking

Who:

Anyone who uses data in Snowflake to understand the contents of the data they are using (data types, PII identification)

Test & Learn:

Created Glue workflows to create data catalogs with 2 different data sources:

Salesforce Account

Data Science (ztrans device data)

1 re-usable script for any data source

R: Led to a x dollar extended engagement with slalom with a team of engineers.

Keywords while being interviewed:

LPs:

Questions it can address to:

8. RFP & Pre-Sales Solutioning

Situation:

We had a few big RFPs come in — some from existing clients and some new logos — and our D&A practice was trying to grow its pipeline for the next fiscal year. I was pulled in as the solution architect to shape the technical proposal.

Task:

My role was to understand the business pain points, come up with a solid AWS-based data solution, and make sure the proposal included a clear architecture, timeline, and delivery plan.

Action:

I worked closely with the client exec, practice leads, and pursuit team to break down the RFP.

I designed the data platform using AWS services like S3, Glue, Redshift, and LakeFormation, and made sure it fit their governance and scaling needs.

I also ran discovery sessions to confirm the assumptions, helped draft the SOW, outlined the RACI and deliverables, and joined the orals to present our solution.

Throughout the process, I made sure the pricing, staffing, and technical scope all tied together cleanly.

Result:

Those efforts contributed to about \$2–3 million in new pipeline for our practice through a few successful RFP wins. It also strengthened our position with the client, leading to follow-on work.

Keywords to use:

RFP cycle • Pre-sales • AWS data architecture • SOW definition • Discovery workshops • Pipeline growth • Solution alignment

Leadership principles (optional tie-ins):

Earn Trust • Dive Deep • Deliver Results • Customer Obsession

9. Partner Collaboration & AWS MAP Funding

Situation:

We had a client who wanted to test a new data architecture, but there wasn't any budget for a proof of concept. We didn't want to lose momentum.

Task:

I needed to find a way to fund the POC so we could move forward and show quick value.

Action:

I reached out to our AWS alliance team and explored MAP funding options.

I put together a short proposal explaining the use case, expected AWS workload, and how it could grow cloud consumption for AWS.

Once funding was approved, I led a small cross-practice team — data and Salesforce — to run the POC and integrate Salesforce Data Cloud with AWS analytics services.

Result:

We got the POC fully funded by AWS, so the client paid nothing but still got a working demo. It built trust, opened up a larger engagement, and strengthened our AWS partnership.

Keywords to use:

AWS MAP • Partner funding • POC enablement • Alliance collaboration • Cost optimization • Cross-practice teamwork

Leadership principles (optional):

Think Big • Invent and Simplify • Ownership • Earn Trust

10. What I want to tell: Building capability gaps through coaching GIC **salesforce** analytics team for the best de practices, was called out by hongwu by customer bridging critical data gaps

S: The GIC analytics team was facing a lot of challenges with data

T: I was brought in as an architect to pener\trate the proposed architecture and best practices in the existing team and

A: I provided detailed technical specs and also had multiple sessions with the client team to help understand the impl level details and the need each layer served and best practices to follow to have a more robust and scalable data platform

R: Adoption of the architecture grew, the lead data engineer was more equipped with taking the data needs with the core eng team, and also have their team follow the standards for implementation, This led to faster and more reliable data insights of teh compensation program. Key client GIC analytics director also acknowledged that I have filled in the critical data gap in the team.

Keywords while being interviewed:

LPs:

Questions it can address to:

11. What I want to tell: **Slalom** CMD5 - partnered with snowflake to place fivetran, dataIQ, thoughtspot, alation and others to define the architecture and present to , placing their value and pitch for customers

S:

T:

A:

R:

Keywords while being interviewed:

LPs:

Questions it can address to:

12. What I want to tell: **Account expansion and extension at fruitco and salesforce**

S: Slalom was finishing the

T: Presenting to the client about data discovery and assessment, and talking through the

A:

R:

Keywords while being interviewed:

LPs:

Questions it can address to:

13. Taking PM hat when the team member was under performing and showing urgency to hire and train the best to meet the needs of customer

14. Constantly learning - with the growth of asking for databricks by customers, I recently did some databricks certifications and also part of partner training. Have recently completed a cohort doing some hands on with rag architecture with proof of concept presenting this our internal group.

15. Performance optimization - I used EMR to optimize which improved query time by 40% and earned trust with the customers using the ICE product

Aprajita -2

Aprajita -2

FruitCo and Salesforce

Co-ordinate with business process people, identifying process gaps, how these business process gaps were filled by technology.

During the discovery we did not only look from the technology side of things, but also looked at the process gaps, how you were **not leading only from the architecture perspective, but how you were leading the business processes**. Where the **business processes working with them** and to identify how these gaps can be overcome with the technology perspective.

(BT team, talking to architects)

While the team was working on the discovery phase, I went ahead and **spoke to the key stakeholders on how they can benefit from ... what are the things you can prioritize to get immediate benefits**. (talk about low hanging/ high hanging fruits) - curated layer first... bringing in governance in curated zone.

How we can integrate with other teams and benefit from

Benefits - quality or quantify.

DG - piled on the last, while you were doing the interviews, you found the process gaps, captured and created the future phase.

They got excited, they wanted to get into the next phase.

I created the opportunity and generated the roadmap, how much TCV? Check from SOW. (2million - 5 year, 3 year deal - 1 million)

BASICALLY - for selling you need to know the numbers, how many people, how much you coordinated, different teams, high level ballpark number is good.

Junior resource - 150 per day

ADVISORY ROLE - phase 1 sell, while the team did wrap up, you extended as an SME.

Taking the vision and strategy alignment with exec leaders.

FRUITCO - leading the whole thing, closely working with key stakeholders to give them a vision, to strategize with them, that generated the prop.

As a leader, you need to show how you are able to co-ordinate

RFP

Why AWS? They were on AWS cloud already. AWS MAP funding.

We are technology agnostic. We helped them evaluate based on their environment, and data volume, ...(for audit purpose they wanted to keep 5 years of dat)

We were getting good funding from AWS MAP funding from AWS for the engagement, that also drove the decision.

Account team gets involved,

IRhythm

In our case, there are multiple account leaders who are involved in getting initial lead, and then the process... we understand the requirements, we give initial proposal, what it will take, it involves solutioning, and trade-offs, and then if they go with us, this is the team and timeline that we propose.

You worked on the logo.

When I joined Slalom, this was my very first project, which was a new logo, and we wanted to make a good impression and build a long term relationship with the customer. As a senior person I was responsible to lead the solutioning, and the problem statement was given to me. - discovery and workshop

Show:

- 1. Working with two deals at a time, also working as SME overseeing few deliverables. Helping out the junior resources if they have any specific challenge.**
- 2. Defining strategy with the leaders at the client side**
- 3. Architecture suggestion, and helping the team overcome challenges as they executed the project.**
- 4. Say if ask about percentile (only if needed) - 60:40. Dont be so specific.**

What were the most difficult technical challenge you faced?

Who has been your most difficult client? How did you convince them?

How do you help your team come out of a problem - technical business?

Trying to win a deal with proper solutioning, highlighting risks and any assumptions made during the process.

Top down solution - high level soln when i dont have many data points

Bottoms up - you have all the data... and then you design

What kind of automation can i bring, how to increase the productivity, make the soln better

Business problem solving

Spound techno functional

Final Exp - Notes

My Experience

[Summary \(Elevator Pitch\)](#)

[1. Delivery clients/projects \(post sales\)](#)

[2. RFP Technical solutioning \(pre-sales\)](#)

[5. Client Workshop: iRHYTHM data platform foundation AWS Workshop](#)

[3. Slalom Alliance Partners led growth \(pre-sales\)](#)

Summary (Elevator Pitch)

Being part of the D&A practice one of the senior solutions architect who knows AWS very well, I was part of multiple proposals, which involves certain....i was creating technical solution documents, engaged in doing workshops with the client to try and understand, collect the data point to do initial assessment, to do initial discovery to make sure my solution aligns with the long term delivery plan. In that phase I have impacted 2-3 million dollars of proposals through RFPs.

My work spans both post and pre-sales — I lead solution design during RFPs and strategy sessions, then guide those same solutions through implementation and adoption. For example, I recently helped a global tech firm build their AI/ML Lakehouse platform to improve data trust and scalability.

1. Delivery clients/projects (post sales)

My Recent teams (reverse order) and customers & brief of what they do, and slalom's role:

a. Salesforce Global incentive Compensation team -

i. Business

The ICM team defines the “how and what” of sales incentives — from compensation plan design to ensuring the correct calculation and accounting of commissions.

GIC Analytics & Ops Reporting Team -Providing data transparency and reporting across incentive compensation — ensuring leaders can track performance, payouts, and compliance.

ii. Business challenges:

The GIC Analytics team's core business challenge is to transform fragmented, complex incentive-compensation data into a single, reliable, and compliant reporting ecosystem that leadership can trust for decision-making and financial governance.

iii. Slalom's Role:

Slalom's Role in the GIC Analytics Engagement

1. Discovery and Current-State Assessment

Slalom was brought in to understand and document the current state of Salesforce's GIC reporting, analytics, and data architecture landscape.

This involved identifying reporting gaps, data dependencies, pain points in the existing SPIFF-driven data model, and inefficiencies in how compensation analytics were being produced and consumed by leadership.

The team worked closely with GIC stakeholders to capture the "as-is" state of reporting workflows, systems, and compliance dependencies (SOX, revenue accounting, etc.).

2. Recommendations and Future-State Vision

Based on findings from discovery, Slalom developed recommendations for the "North Star" future state of GIC Analytics — defining what a modernized, reliable, and scalable reporting ecosystem should look like.

This included data architecture recommendations (how SPIFF data should flow into analytical schemas), reporting layer redesign, and process improvements to reduce manual intervention and ensure audit-ready data.

3. Collaboration with BT and GIC Teams

Slalom's role extended to aligning with the internal BT (Business Technology) team to ensure that the recommended analytics and architecture roadmap fit within Salesforce's enterprise data strategy.

They acted as a bridge between the GIC business teams (who defined the reporting needs) and BT technical teams (who owned data systems and integration pipelines).

The goal was to jointly define a North Star vision and architecture — a unified approach for data ingestion, transformation, and reporting that could scale across all 17+ compensation plans and ensure financial compliance.

iv. My Role (Elevator Pitch):

I led the discovery and architecture strategy for Salesforce's Global Incentive Compensation (GIC) Analytics transformation — a cross-functional initiative aimed at unifying how Salesforce manages, calculates, and reports on global sales compensation.

My role centered on conducting a current-state assessment across the SPIFF platform, data pipelines, and reporting layers to identify architectural and reporting gaps. I partnered with GIC business leads and the BT (Business Technology) team to define the "North Star" data and reporting architecture, ensuring the new ecosystem could deliver SOX-compliant, scalable, and automated analytics across 17+ incentive plans and thousands of sellers.

In short, I helped Salesforce move from manual, fragmented reporting to a centralized, governed, and architecture-aligned incentive analytics platform that leadership can trust for performance and financial decisions.

b. Apple - Global Sales Team -

i. Business:

Apple Sales Operations – Business Overview

Mission & Scope

Apple's Global Sales Operations team, particularly Ben's US Channel group, is responsible for enabling Apple's consumer and carrier sales partners (Amazon, Best Buy, Target, Costco, Walmart, AT&T, Verizon, T-Mobile, etc.) to drive revenue growth through data-driven insights, tools, and operational excellence.

Core Responsibilities

Performance Tracking & Reviews: Deliver weekly, monthly, and quarterly sales reviews to monitor channel performance, highlight trends, and identify sales opportunities.

Partner & Marketplace Insights: Provide visibility into how Apple products are performing across major online and retail channels.

Digital Sales Enablement: Expand capabilities in digital sales analytics — understanding traffic, conversion, and deal performance across digital storefronts.

Operational Support: Equip field sales teams with dashboards, scorecards, and tools (e.g., Tableau, Salesforce D360 Digital, Rocket) to understand short-term performance and inform sales strategies.

Governance & Data Access: Manage data ownership, access control, and permissions across teams and tools.

Business Challenges

1. Data Fragmentation & Inconsistency

Data is sourced from multiple partners and third-party providers (Amazon, Best Buy, SimilarWeb, etc.), each with different formats, time scales, and update cadences.

Much of the data is manually pulled or partially structured, making it difficult to harmonize into a unified dataset.

Impact: Incomplete and inconsistent data undermines trust in reports and limits Apple's ability to make timely, data-driven sales decisions.

2. Lack of Harmonized Digital View

The team's "North Star" goal is to build a cohesive, weekly digital sales dashboard showing traffic, search trends, and sales across all partners.

However, achieving this requires standardized ingestion, cleaning, and transformation pipelines that don't yet exist.

Impact: Apple lacks a single source of truth for digital sales performance — making it hard to link web traffic and search data to actual sales outcomes.

3. Limited Tool Adoption & Data Trust

Sales teams underutilize tools like Tableau, Salesforce D360 Digital, and Rocket, often due to complex dashboards and low confidence in the underlying data quality.

Without consistent adoption, insights fail to translate into action.

Impact: Operational inefficiency and missed opportunities for proactive sales management.

4. Insights vs. Reporting Gap

Current focus remains on report production, while the long-term vision is to deliver actionable insights that drive decisions (“what to do next,” not just “what happened”).

There’s also a need to improve data literacy and build a culture of self-service analytics among field teams.

Impact: Sales operations remain reactive rather than strategic, limiting Apple’s agility in a fast-moving retail and carrier landscape.

5. Alignment & Governance Complexity

Organizational alignment is needed across consumer, carrier, and digital sales operations, with stakeholders owning key parts of the process.

Data ownership and access governance are bespoke and siloed, slowing down standardization and cross-channel visibility.

Impact: Slow progress toward a unified data strategy and difficulty scaling analytics capabilities across regions.

ii. Slalom’s Role

Conduct a comprehensive current-state assessment of Apple’s Sales Operations data ecosystem, processes, and reporting tools.

Engage stakeholders across Consumer and Carrier sales functions to capture key business requirements, KPIs, and success metrics.

Define the future-state architecture for unified digital sales reporting — outlining both short-term quick wins and long-term strategic recommendations.

Enable and educate Apple’s internal engineering and analytics teams on the proposed architecture, data harmonization strategy, and implementation roadmap.

iii. My Role (Elevator Pitch – Draft)

I led the data and architecture strategy for Apple's Sales Operations transformation — focusing on harmonizing fragmented digital sales data across major retail and carrier partners.

Through stakeholder interviews and technical assessments, I defined the future-state reporting architecture and established a roadmap to transition from manual, partner-specific reporting to a standardized, insight-driven analytics platform.

My role also included guiding Apple's internal engineering team on data integration best practices, governance, and scalable implementation patterns to ensure lasting adoption and measurable business impact.

c. Apple - AIML Annotations Operations team



Business Overview

Apple's AI & Machine Learning (AIML) organization powers intelligent features across core applications such as Siri, Mail, Calendar, Music, and Photos.

Within this group, the Annotation Operations (Annotation Ops) team provides the human-in-the-loop foundation for Apple's machine learning — a 6,000-person global workforce responsible for labeling and validating data that trains Apple's ML models.

Key Business Functions

Human-Guided Machine Learning: Analysts annotate text, images, audio, and video to improve model performance and accuracy.

Massive Operational Scale: Hundreds of millions of annotation records are produced daily across multiple global sites.

Structured Workforce Model: Multi-tier hierarchy (Analyst → Team Lead → Site Lead) operating under strict throughput and quality metrics.

Data-Driven Operations: Productivity, quality, and quota compliance are tracked continuously through dashboards and performance analytics.

Slalom's Role: Building data pipelines, curated datasets, and high-performance Tableau dashboards to monitor quality, throughput, and operational efficiency across this massive ecosystem.

- i. Slalom's Role:

6. Slalom's Role & Apple's Environment

Slalom's role has evolved from data visualization (Tableau) to data engineering, curation, and performance optimization.

Apple's AWS and Snowflake environments are customized, highly secure, and integrated with Apple's private infrastructure — not standard cloud setups.

Emphasis on simplifying data consumption for Apple teams while maintaining trust in data quality.

- ii. My Role (Elevator Pitch):

Work with the

d. Genentech (Jan 23) -

- i. Business:

the company's Research and Early Development (gRED) group is responsible for discovering and developing new medicines. Its work involves fundamental scientific research, tackling complex diseases like cancer and multiple sclerosis, and using advanced computational tools and data science to create breakthrough therapies.

gRED scientists need a cloud-enabled computing environment that enables seamless collaboration with our external collaborator, Recursion, and our Roche colleagues in pRED. They wish for computational scale on demand with data, methods and infrastructure sharing that will speed the generation of insights and the identification of novel targets and small

molecules. gRED and Recursion will work together in partnership to achieve this vision.

Most of gRED's data is on premises and fragmented. This aspect of the current state of gRED's data landscape could slow down cloud adoption and should be factored in when measuring the pace of adoption. This risk is likely to materialize in scenarios where new applications scheduled to be deployed to the cloud need to integrate with current on premises datasets. Deployment of such applications may require additional design considerations to address security and performance challenges related to the transfer of data between on-premises and the cloud. This is particularly true for any new cloud application which would need to integrate with gRED's large Oracle instances such as TAPIR, BioPRod, ACCPRD2 or SSP or datasets located on NFS and GPFS.

Finally, while gRED users have expressed an interest in improved dataset management practices, a culture of systematic data management is currently not established within gRED. Ideally, gRED's cloud offering should provide solutions to help users adopt these practices and develop such a culture and keep its cloud as a clean and organized space from the data standpoint. This could pay dividends once a large number of datasets and applications are deployed and it is possible to list out most datasets along with their metadata.

- ii. Slalom's Role:
- iii. My Role (Elevator Pitch):

e. IRhythm (oct 22 - dec 22)-

- i. Business -Based on the sources, the company, referred to as "rhythm" (and the subject of the audio file "Irhythm.mp3"), appears to be involved in heart-related technology.

In simple terms, the company is described as being like a "Samsung smart watch or Apple Eyewatch but specifically made for heart".

The sources also detail the current scope of the consulting engagement with the company, which focuses on their data infrastructure and strategy:

- Data Strategy: The current engagement is a strategy engagement that involves developing a data strategy for services for the self-serve sort of environment.
- Infrastructure: The company uses AWS. They have decided to move from Red Shift to Snowflake.

- Data Handling: The engagement involves work on data injection. Currently, they are writing their own extractors in Golang because they do not have a full-fledged ETL or ELT tool yet.

- Project Goals: A specific job during this engagement is to build an architecture and a framework where people can onboard data sources in a consistent fashion onto the data lake. There is also a mention of reading about Argo workflows for orchestrating containers.

- ii. Slalom's role - new logo for slalom team, new client. Slalom was in 3rd week of strategy engagement - a lot of moving parts, they don't have full fledged ETL or ELT tool , they were writing their GoLang extractor, they have decided to move from redshift to snowflake - that aspect also. We are doing data strategy for data services - catalog service. Build an architecture where people for ...workflow orchestration. 4 phase explore, design and develop, do a test and learn experiment, define a roadmap


f. **ICE Mortgage Data Connect -**

- i. Business:

Data Connect

ICE Mortgage's **Data Connect** transforms complex, high-volume Encompass loan data into **structured, queryable datasets** for lenders — but faces challenges around **real-time data consistency, schema evolution, massive scale, and multi-tenant orchestration** while maintaining performance and reliability.

Components



- ii. DE Team's Role:
- iii. My Role (Elevator Pitch):

Leading the data connect team, for building the dta aconnect products, working with product managers on requirements and feature releases, and also work with AWS Sa for architecting and designing the data lake.

2.RFP Technical solutioning (pre-sales)

Elevator Pitch:

Being part of the D&A practice one of the senior solutions architect who knows AWS very well, I was part of multiple proposals, which involves certain....i was creating technical solution documents, engaged in doing workshops with the client to try and understand, collect the data point to do initial assessment, to do initial discovery to make sure my solution aligns with the long term delivery plan. In that phase I have impacted 2-3 million dollars of proposals through RFPs.

How did you do it - i worked with the pod, i worked with the account team, i worked with RFP timelines, depending on what their requirements is, we responded in a certain way,

I have also helped with the SOW, writing down the activities, what deliverables, how the RACI will be defined, what are the technical components or automation, which should be involved in the solution.


Impact - it generates the pipelines for our D&A, brings business to the company, helping Slalom grow.

RFP Responses - Examples

1. Informatica Integration Planning Program

- a. Client: Salesforce
- b. Date: Aug-2025
- c. Description: The acquisition of Informatica is currently under regulatory review and is expected to close in Q1 FY27. Integration planning activities have begun pre-Close to get ahead of impending scale & complexity of the deal. The Salesforce Customer Success org is looking for consultants to help us with pre-Close preparatory work with the goal of faster and smoother integration post-Close.
Scope of work details:
Note: Below scope applies to all functions under the ‘Customer Experience’ organization at Informatica viz. Customer Support, Customer Success, Renewals, Professional Services (Education Services only) and Strategy & Operations:
Business Process Analysis & Customer Journey Mapping
Systems and Data Flow Assessment
Employee Role Mapping
- d. My Role: Contributed to Systems & data : Approach and Deep Dive slides and presentation for below slides:

Systems & Data: Approach



Systems & Data: Deep Dive

Identify systems used across Customer Experience, understand use cases, customization levels, and ownership to inform integration planning.

Area	Acceleration Methods	Key Questions	Assessment Criteria
 Systems	- Leverage Slalom's Salesforce-on-Salesforce architecture knowledge - Leverage pre-built system inventory templates for faster discovery	- What in-flight or planned migrations could impact timelines? - Which upstream dependencies are most critical or fragile? - Which customizations or non-standard configs could complicate migration?	- Alignment to Salesforce core and integration priorities - Migration readiness, including in-flight or planned changes - Ownership clarity and customization complexity
 Integrations	- Use mapping accelerators to capture integration points & dependencies - Apply upstream/downstream tracing methods to surface hidden risks - Leverage integration architecture playbooks for quick pattern identification	- What integration architecture patterns are currently in place, and how scalable are they? - What dependencies and critical flows exist? - How are errors handled to prevent failures from cascading across systems?	- Resiliency of connections and error recovery - Impact of upstream/downstream dependencies - Fit of architecture patterns for future state
 Data	- Apply Data Quality & Readiness Framework for migration-ready datasets - Apply scoring models to prioritize migration readiness - Leverage IDMC & Salesforce data model expertise to align structures & flows	- Where is data quality weakest or least documented? - What compliance or residency rules could restrict movement? - What planned updates to data models, mappings, or transformation rules could affect migration or integration?	- Data quality, completeness, and core object alignment - Business criticality to CX processes - Readiness for migration, factoring model or mapping changes

slalom ©2025 Slalom. All Rights Reserved. Proprietary and Confidential.

15

2. Data Aggregation and Reporting Tool RFQ – Energy-RFP

- a. Metropolitan Transportation Commission
- b. March-2024
- c. Description: BayREN currently administers a portfolio of six energy and water efficiency programs. Each program has a designated Program Manager/Lead who maintains a separate database housing their program's data. A central database/repository does not exist; instead, there are currently six isolated datasets requiring manual, time-intensive processes to display the impacts of multiple programs (i.e., multiple databases). This project seeks to create a Data Aggregation and Reporting Tool that reduces this burden by automatically interacting with the six existing databases to retrieve data on a predetermined frequency (e.g., monthly), transform the data into common metrics, and consolidate data from multiple programs into customizable reports.
The resulting tool should be hosted by ABAG/MTC and accessible to approximately 25 users, including staff from ABAG/MTC and each of the nine counties. It is desirable that permissions be structured to permit a small subset of users with the ability to create and edit reports, while the majority of users would only be able to run preset reports.
- d. My role: Writing technical specs for the RFP response for following sections:
 - i. Approach for the Technical Solution
 - ii. Example of a Cloud Hosted Data Platform Using Amazon Web Services (AWS)
 - iii. Assumptions / Constraints in Selecting the Technical Solution

The components highlighted in blue are the ones that would be implemented in the Data Aggregation and Reporting tool based on our present understanding. The design of the tool is extensible to include the features depicted in the overall diagram.




Figure 2 - Standard Architecture for the Proposed Solution

Links:

[Solutions, Insights and Experiences - Slalom](#)

<https://lucid.co/blog/mckinseys-three-horizons-of-growth>

3. data strategy workshop proposal (CCC)

4.

5. Client Workshop: iRHYTHM data platform foundation AWS Workshop

Project Description: Client is interested in exploring options for establishing data analytics platform and jump start implementation of high priority use cases. Client needs to define a roadmap for its future data and analytics capabilities which potentially require AWS data analytics & machine Learning capabilities and services. With the focus on

A- Basics functional and non-functional needs for establishing an MVP for the data foundation platform.

B- Ability to collect data from multiple sources and ingest into the analytics workspace

- C- Build a roadmap for building a foundational data analytic capabilities/service that enables the client to create self-service BI and analytics ecosystem for the business users.
- D- Analyze requirements for data quality and governance requirements

Slalom / AWS Outcome	client on their data cloud journey roadmap, starting point and the MVP of the data . el requirements for data collection, storage, governance and analytics solution and ll leveraging AWS services. el architecture/component for the target ecosystem
Customer Outcome	the starting stages of the client's data transformation journey mainly driven based uture need for leveraging specific data analytics and machine learning services and ies that are required for MVP of the data platform and learning about the AWS es and funding the program and reduces the initial cost of the business.

3. Slalom Alliance Partners led growth (pre-sales)

I also have experience working with our alliance partners within Slalom, in which I have worked to get some MAP funding, or **funding** from technology partners, which we have invested in doing small POCs, or trying out use cases at the client side.

You also worked with multiple practices. I have worked with Salesforce practice, to talk about the **salesforce data cloud integration**,

1. Brickbuilder Specialization
2. Contribution in CMDS
3. AWS MAP

16. SOW Building

=====

3. Discovery of data and current state and putting lens of gaps compared to best practices
4. Evaluating and Presenting architectural options
5. Technical advisor for data architecture and governance
6. Build SOW
7. Provide technical solution in RFP responses
8. Workshops with customer

9. Stakeholder interviews and building requirements specs
10. Partner growth and funding
11. You **work alongside account executives or partner managers** before the deal closes
12. You **understand client needs, design technical solutions, build prototypes/demos,** and **present architecture diagrams** that lead to adoption.
13. You **translate business pain points into technical proposals** and **influence buying decisions.**
14. You **enable partners** — helping them co-sell and co-deliver .

companies-business

ARCHIVED - NOTES

[Index](#)

[a. Salesforce Global incentive Compensation team](#)

Index

a. Salesforce Global incentive Compensation team

- Incentive Comp Plans: Salesforce utilizes incentive comp plans. It is believed that the company has 17 different comp plans and incentive comp plans.
- Commissions Logic Architecture: The SPIFF Proserve team is responsible for the commission's logic architecture. This team's job is to set up the commissions logic so that the end result—the comp statement delivered to every seller—is accurate.
- Compensation and Revenue: Commissions are based on a sold opportunity or ACV (Annual Contract Value). The Accounting and Payile Process (led by Dan Riley within GIC) handles the accounting for these commissions, ensuring compliance with revenue recognition rules, as a sold opportunity generates revenue over time rather than immediately.

The GIC (Global Incentive Compensation) team owns several workstreams, one of which is Analytics and Ops Reporting. This workstream is led by **The GIC (Global Incentive Compensation) team owns several workstreams, one of which is Analytics and Ops Reporting. This workstream is led by Tim Lynn.

Based on the sources, the primary responsibilities and functions of the GIC Analytics and Ops Reporting workstream include:

- Building Reports: The people involved in this area (the two individuals being addressed) will "build the reports that the leaders will see" and "figure out what needs to be done" based on those reports.
- Servicing Leaders' Reporting Needs: The end goal of the overall program is for a group of sales leaders and finance leaders to "receive

these daily reports, monthly reports, quarterly reports, all the different report types". The work done by Tim Lynn and the team has to be "serviced or... rendered by the BT team" to deliver these reports accurately and seamlessly.

- Addressing Reporting Gaps: A critical starting point for the people joining this workstream is to determine the "reporting gaps piece of it" and figure out "where is the state of affairs when it comes to reporting needs and gaps".

The overall context of the GIC's work, which includes analytics, is focused on the accuracy and delivery of compensation-related information:

- Data Architecture Dependency: The ability to deliver these reports is heavily dependent on the data architecture. Mark, the GIC leader, brought in external support due to challenges with the architecture piece. The analytics and reporting work ties into understanding how the data from SPIFF is stored into the schema, tables, and made "ready for all the reporting".
- Compliance and Financial Reporting: Some of the reports generated must be Sox compliant and are intended to "feed the financial flow statements", meaning there is "no room for error".
- Leadership Oversight: The team is expected to meet with Tim Lynn to "nail down" the state of reporting needs and gaps. The expectation is that they will look at existing work to understand what has been done regarding reporting gaps.

In summary, GIC analytics is the team responsible for building and ensuring the accurate delivery of compensation and operations reports to sales and finance leaders, a task that currently requires significant focus on addressing reporting gaps and underlying data architecture challenges

b. Apple - Global Sales Team -

Team Focus & Responsibilities:

- Ben's team is focused on the US channel, covering both consumer and carrier reseller partners. This includes major retailers like Amazon, Best Buy, Target, Costco, and Walmart, as well as carriers like AT&T, Verizon, and T-Mobile.
- The team supports sales operations by providing tools, insights, and strategy to help sales teams understand marketplace dynamics and improve sales performance.

Sales Operations and Functions:

- The team is responsible for preparing weekly, monthly, and quarterly sales reviews, helping sales teams understand short-term performance (current quarter + one quarter) and identifying trends and areas of opportunity.
- Ben emphasized that digital sales operations are relatively new for his team, and they are still building up their capabilities in that space.
- A key responsibility of sales operations is helping sales teams analyze the performance of specific sales deals and understanding why certain deals may have underperformed.

Challenges with Data:

The team faces significant challenges with the inconsistency and disparate nature of the data they receive. Data comes from various sources (e.g., partners like Amazon, Best Buy, third-party providers like SimilarWeb), and it's often in different formats, on different time scales, and sometimes needs to be manually pulled.

Harmonizing the data is a major goal, with the aim of building a cohesive view of digital sales across all partners. This includes correlating sales data with digital traffic and understanding the factors driving sales.

Key Project Goals:

The "North Star" for the project is to be able to report digital metrics (sales, traffic, search trends) across all partners in the US channel on a weekly basis.

There is a strong need to develop standardized processes for ingesting and harmonizing the data, which is currently incomplete and unstructured.

Success will be measured by how well the sales teams adopt and rely on these insights.

A key indicator of success is if sales teams use the tools regularly and can take actionable insights from the data.

Tools & Data Consumption:

The team uses several tools for data visualization and reporting, including Tableau and Salesforce (D360 Digital, which integrates data from various sources). There is also a tool called Rocket for sales teams.

One challenge is that the existing tools and dashboards are underutilized by the sales teams, often due to a lack of trust in the data or complexity in the reports.

Long-Term Vision:

The goal is to transition from merely producing reports to providing actionable insights that can drive sales decisions. Ben emphasized that insights should suggest actions and not just present raw data.

Adoption of these tools by the sales teams, with superusers acting as champions for the data, is critical to success. The project aims to improve data literacy and encourage the use of self-service reporting tools.

Organizational Alignment:

Ben sees the project as building a roadmap to harmonize data and solve the challenges around sales operations. He believes that creating a detailed plan over the next eight weeks will provide a strong foundation for future work.

He mentioned key stakeholders like Joel and Jigger, who are leaders in the US channel sales organization, as well as the need for alignment across the sales operations and consumer/carrier sales functions.

Data Ownership & Governance:

Ownership and access to the data are managed by the sales operations team, with permissions often being bespoke to the needs of the carrier and consumer teams.

Salesforce and Tableau handle the technical provisioning of access.

Apple AIML team:

i. Business:

The discussion explains how Apple's AIML (AI & Machine Learning) organization — originally focused on Siri — now supports a wide range of intelligent features across apps like Mail, Calendar, Music, and Photos. The focus is on the Annotation Operations (Annotation Ops) team, their structure, data scale, and the role Slalom plays in supporting data analytics and platform performance.

1. AIML & Annotation Ops Structure

AIML Group: Handles all ML-enabled features (voice, text, image, etc.) across Apple products.

Annotation Ops: A 6,000-person global workforce providing “guided machine learning” — humans labeling data (e.g., drawing boxes around objects, transcribing audio, translating text) to train and validate models.

Tasks: Range from labeling images and video frames to audio transcription and translation.

Example: tagging whether a music feature correctly removes vocals or identifying text in multiple languages.

Nature of work:

Highly repetitive and measurable — throughput and quality are key metrics.

Workers are often early in their careers; high churn due to monotony.

Teams operate like a factory assembly line with quotas and performance dashboards.

2. Organizational Hierarchy

Regions: Americas, Europe, and Asia-Pacific.

Each region has regional leads and multiple sites.

Each site has:

Analysts → Team Leads → Senior Team Leads → Site Leads.

Analysts are grouped by locale/language expertise (e.g., English, Chinese, Japanese).

Sites are selected for multilingual capacity and data security — most work is done on-site, not remotely.

3. Workflows & Quality Measurement

Tasks are broken into projects requiring specific quantities of labeled data (e.g., “10,000 data points in a week”).

Performance is tracked via throughput and accuracy metrics.

Outlier detection identifies workers who may be underperforming or gaming the system.

Analysts can “level up” to more complex tasks or QA roles if they consistently perform well.

Those who drop in quality are retrained or reassigned.

4. Data, Scale, and Tooling

The data volume is massive — hundreds of millions of records from annotation tasks logged daily across thousands of workers.

Slalom supports Apple in building data pipelines, curation layers, and Tableau dashboards to track performance and quality metrics.

Performance tuning is essential — dashboards must handle massive datasets efficiently at the database layer.

5. Key Stakeholders & Roles

Deborah (Apple): Senior leader (29 years at Apple) overseeing annotation ops analytics; values Slalom's data partnership.

Bryce (Slalom): Focused on Tableau performance and design optimization.

Chad / Karen / Sean / Tully: Slalom team members supporting different parts of the analytics pipeline and transition planning.

7. Key Takeaways

Apple's AIML annotation system is industrial-scale human-in-the-loop ML — labeling, QA, and model validation.

Organizational complexity reflects language diversity, massive scale, and sensitivity of data.

Slalom's challenge is building scalable, performant, and trustworthy analytics to help Apple leaders understand and improve global annotation operations.

Culturally, many analysts aim for “average” performance to avoid scrutiny, creating an analytical need for outlier detection and motivation frameworks.

ICE Mortgage (Data Connect)

The loan data gets originated on Encompass. Customers want that loan data. They can get making API calls to getLoan, they can use the SDK, they could go through developer connect, but they would have to get the loan, parse it and do some work to get that data.

Sometime back, why don't we get that data, break it apart and save it to database, and give the jdbc uri endpoint, so that they can pull it. It started off as SQL, also supported by Postgres.

Gist is you get the endpoint, and the data would reside in the database. That database would have around 200 tables, a lot of columns, and the way it works, we internally take that json, that json has nested structure, so we say that these child would go in table 1, these thing going to table 2, another many-many, in other words, every section of the loan is translated automatically to a table, and all the properties inside of these map to column.

We get the loan, we do the processing and split the loan, and we insert it into database, and we keep only the latest version of the loan in db, in other words if you see loan guid 12345, you are going to see one row for that loan, in all these tables, some of these tables might have

You are not going to see loan 12345 at 1 pm, loan 12345 1:05, you are not going to see multiple rows, you are going to see only one row, and the timestamp that happens is the latest one, or the most recent one. We update the older version, we do a lot of updates

A loan changes around 90/80 times per minute or per hour. Multiple loans are worked on, we insert and update those many loans.

EBS publish the loan change event writing to rabbit mq. We have one component called LQI - loan queue ingestor, that guy takes it from rabbit mq and passes it to Kinesis. It does not do much processing, it is kind of bridge, from this Kinesis the loan queue processor (LQP) kicks that loan id that is there in that message (the full loan is not there, it is just a wakeup call that this loan changed). The payload of the event has basic meta data, (loan id, timestamp, clientid, instanceid), with that info this guy goes and does the GetLoan call. Once it gets the loan, it sends the full payload of the loan to the Kinesis.

From there, there are 2 branches,

one branch goes to the lambda archiver - it compresses and encrypts that loan and it saves it in the Global data lake. Global data lake is nothing but S3, with some sort of a partitioning strategy, like a file system path structure, clientid, instanceid, year, month, days, very common way of storing data in S3. People talk about it as partitioning strategy, why is it partitioning, because you have effectively segregated the way you are saving data, and you are partitioning your data, you are slicing your data by client id, instanceid, year, month, date, that is like a bucket, all the loans that match this go into this bucket. There would be many loans matching this strategy, many loans are saved on a daily basis. The partition strategy kind of dictates how large is gonna be the number of loans on those buckets meaning how many loans are saved per day, imagine that your partition strategy was not up to the date, but upto a year - that would be too many.

As a data engineer, you need to think about how you partition and slice your data so you can process it in manageable chunks, how do you divide

and conquer this terabytes of data, you cannot, you have to have partitioning instead of buckets.

So we store this data here, here we have JSON compressed and encrypted,

and we also send these payloads to the forwarder lambda, (confusion - this guy encrypts the data), it can be encrypted here or here.

Lambda forwarder also picks up from Kinesis, and it is called forwarder, because each data connect customer has his own swimming lane, independent path with infrastructure that belongs to that customer, that process the data for that customer only, so this is what they call PDL Kinesis, this is a Kinesis stream only for that customer, its independent from these other Kinesis.

And then there is a loan processor, this is an ECS containerized application, this is a Java application, this is the one that does the Json split into tables, it generates a bunch of INsert/update commands, and then goes to the database and says execute this batch.

Then there is a RDS instance, SQL server or Postgres, and the forwarder is kind of a router, it sees one loan and it has like basic metadata like clientid, instanceid, where should I send you, oh based on metadata, lookups table, it knows it has to send it to this customer, that routing is pretty much the work that it does,

Loan processor picks up does splitting, does insertion, and this is what kind of data connect is.

The output to the customer is the database with loan data. The way they use the database is usually by select * from a loan where the modified timestamp is later than the last timestamp i queried. I remember the last query time (say 3 o clock), next time i query is tell me anything that changed after 3 o clock. Sometimes they could be running queries against 200 tables, or they may be running against specific tables, we don't control that,

We keep the latest version of the loan, but that version keeps on updating modified, timestamp.

There is another microservice which is not highlighted here which is called DPS - data platform service. What we just saw is a real time streaming pipeline, but data connect is not all about this. One question is how are these things created, or provisioned, well all these things are created dynamically, automatically, when there is a customer onboard, for which we need another component called DPS.

DPS provides a rest endpoint, and the customer signing up for data connect, that customer was already provisioned on Salesforce, somebody on Salesforce said you have access to data connect and I am going to give you a skew or entitlement that authorizes you to use that, so when you login into that UI, we go to PSS there is a ... there, that gives you a token, that entitlement tells that you are a data connect person, in that you have form, in that form you say that I want to access my database, with username and password. At that time, the database did not exist. When they fill the form, DPS starts provisioning, go and create the PDL kinesis, go and create the ECS loan processor, go and create the RDS instance for this customer.

Second thing that happens is you need to create the older tables, that needs to be there, for you to start inserting data,

The flow we talked about is about inserting data, but who created the tables for us to insert the data, DPS creates the table. How does it know that it needs to create 200 tables or 201 or 203 and these many columns.

Encompass has a loan schema, description of 19.3, 19.4, 20.1. Every single release of Encompass has a schema. These guys look at the schema, and they translate that into a bunch of tables. That's the second step, first step is we provision all these things, then we set up the tables, now we need to replay your last three years worth of data and put them into database, because if you only do this real time thing, you are going to start seeing fresh data as it happened right now, you are not only interested in that, you are also interested in historical data, so we need to do what we call data replay.

Data replay is a separate service, it exposes an API. Data replay is orchestrated by DPS, so DPS calls replays service, it says replay the last three years of data, it takes a data range. This guy goes to EBS, fetches the loan, sends it to Kinesis again, and from that Kinesis, it starts flowing. Data lake gets refreshed as a side effect, but the point is that it uses the same kind of infrastructure, so that is good and is bad. Sometimes we take a week's worth of data, and replay to it. That happens when we onboard a customer, as well as everytime the customer switches a new Encompass version, 20.1, 201. For multiple reasons, the schema might change, that may require us to replay historical data to make it look exactly like your latest version of data. So we do a data replay effectively every quarter for every customer.

Hot swap mechanism for these kinds of scenarios.

Historical data get stored in RDS, at the same time new loans coming up, every time you save a loan, you have a new version on the loan, it's like an index that keeps on increasing and you also have modified UTC timestamp, Data Connect uses that to prevent overriding fresh data with old data.

Let's say that in this pipeline you have the latest version of this loan, you are also doing data replay and so how you got an older version of that loan, you are going to lose the fresh data that came first? NO - because you are going to check what version it is, when was it produced/modified. Oh this thing, this is older, so I am not gonna update it. So it takes off not inserting old data.

There is also a hash involved. And the purpose of the hash is following - you have a new json document every single time, not everything on this json changes, it turns out that maybe 80% of the loan did not change, only 20% of the loan changed, so the way they manage that, they do not want to do a full insert/update on the whole thing, because that would be a waste, so for these things, they calculate for every single loan, of these json document, they calculate hash, I believe we are using MD5, and they compare the hash, if nothing changed, hash is same, I am not gonna do an update on these tables.

Two mechanisms - one is the version and the timestamp, to prevent updating information that is older, and the second mechanism is the hash to prevent updating things which have not changed, so those are two things.

We talked about -

provisioning,

Creating the tables

Data replay

Real time streaming

The only thing that is missing is hot swap.

Hot swap - initially without hot swap, what's gonna happen is data replay might take a week so imagine that your schema changed, you now have three more tables, or maybe something changed, and we need to replay your data for last 3 years, but you also use your database, so we are now loading that bunch of data with new columns with something and that might affect your scripts, might slow down your queries, it is disruptive.

What these guys did was why don't we just create second RDS instance, we populate everything here, according to the new schema, and once it is up to date, and we ask, customer do you want to switch from 19.2 to 19.3 - we switch the dns endpoint, and now we point to the new schema, and the old database, is there around for a week or so, and we say you have it, I am going to tear it down, I am not going to pay for the db... its some sort of blue green deployment, that is faster than replay and it is less intrusive, and you are not affected meanwhile while I am doing data replay.

Data replay happens during onboarding and it happens every quarter, but again think about data replay as a building block, data replay for what when. It might happen for different reasons, but data replay refers to the process of going to EBS, get the loan, replay it through here, so it flows here,

On the second database, on the hot swap, we do both....

Loan Audit - as a product its dc.... Athena spark, is more challenging.

Break the components

What are the challenges?

Interviewing Techniques

Interviewing Techniques

Question Type	Target Time	What Interviewers Expect
Behavioral (STAR)	2–3 minutes	Full story with Situation, Task, Action, Result — concise but complete.
Technical Design / Scenario	4–6 minutes	Talk through approach, trade-offs, architecture diagram (verbally), and reasoning.
Follow-ups / Deep dives	1–2 minutes	Quick, direct answers with 1 supporting detail.
“Tell me about yourself” intro	1.5–2 minutes	Enough to show trajectory, impact, and relevance without rambling.
Closing questions (“Why Databricks?” “Any questions for us?”)	1 minute each	Strategic, thoughtful, and high-level.

Behavioral Question (2–3 minutes total)

Example: *“Tell me about a time you led a difficult project.”*

Here’s what usually happens:

1. **You talk for 90 seconds** → set up Situation, Task, and part of Action.
2. **Interviewer interrupts:** “Why did you take that approach?” or “How did you handle conflict?”
3. You respond briefly (20–30 sec).
4. **You finish** with the Result and reflection (30–45 sec).

 So it’s still your 2–3 minutes of *speaking time*, but over a 4–5 minute total interaction.

How to Fit Technical Detail Into 2–3 Minutes

Use this time structure (the “hybrid STAR + Design” model I coach candidates on):

Section	Time	What to Focus On	Example
Situation (30 sec)	Set technical & business context	“The client’s streaming pipeline was dropping messages due to inconsistent schema evolution.”	
Task (30 sec)	Clarify your ownership & challenge	“My role was to stabilize ingestion and redesign the Lakehouse pattern without disrupting production jobs.”	
Action (90 sec)	<i>Blend technical and leadership narrative</i>	“I redesigned the schema validation layer using Auto Loader’s cloudFiles schema evolution... involved multiple teams, presented trade-offs between append vs merge, and gained client sign-off on phased rollout.”	
Result (30 sec)	Outcome + learning	“We reduced ingestion failures by 80%, cut pipeline latency by half, and built team alignment on a scalable framework.”	

→ Total **2.5 minutes**

→ Enough to show technical depth + leadership impact

Technical Design Question (4–6 minutes speaking)

Example: “How would you design a scalable ingestion framework for streaming and batch data?”

Breakdown:

1. **Minute 1:** Clarify scope & ask smart questions (show structured thinking).
2. **Minutes 2–3:** Lay out high-level architecture (talk through ingestion → processing → storage → governance).
3. **Minutes 4–5:** Deep dive into trade-offs (e.g., Spark Streaming vs Auto Loader, Delta caching, cluster types).
4. **Minute 6:** Wrap up with results, cost/performance trade-offs, or possible extensions.

💬 The interviewer may jump in — that's *good*.
You want a **conversational design session**, not a monologue.

Best Practice: Build Natural Pause Points

Add soft breaks where the interviewer can engage, like:

“At this point, I’d validate if the data latency requirement fits this design — would you like me to expand on that part?”

This keeps it interactive and lets you **control pacing gracefully**.

OpenAI partnership

Scope of Work: OpenAI Codex Pilot

What's happening:

Slalom is participating in an **early pilot with OpenAI's Codex**, an AI-powered code composition and agentic development tool. A small internal cohort (including you) will be granted Codex licenses to **experiment, provide structured feedback, and shape Slalom's long-term adoption strategy** for generative AI in software and data engineering.

Core objectives:

1. **Hands-on experimentation:** Use Codex for internal or personal projects to understand its capabilities in real-world development workflows.
2. **Insight generation:** Document strengths, gaps, productivity impact, and edge cases across different languages and engineering roles.
3. **Operational readiness:** Develop early thinking on playbooks, best practices, and governance for future enterprise rollout.
4. **Strategic assessment:** Evaluate ROI, delivery acceleration potential, and developer comfort levels to inform a broader enablement plan.

Guardrails:

- Use only for **non-client** or explicitly approved work.
 - No client data; focus on internal experimentation.
 - Expect to share qualitative and quantitative feedback after trial use.
-

Strategic Impact on Slalom

1. Workforce Enablement & AI Fluency

- Builds AI-native habits among consultants and architects, strengthening Slalom's reputation as an early adopter of generative development.

- Equips delivery teams to intelligently use and govern AI assistants in engineering and analytics contexts.

2. IP & Methodology Development

- Generates internal IP (prompt libraries, reusable templates, workflow standards) that can be productized into internal accelerators.
- Positions Slalom to standardize AI-augmented development workflows for future internal and client projects.

3. Business & Delivery Efficiency

- Early data on productivity lift (e.g., code generation time, defect reduction, onboarding speed) will inform whether Codex adoption drives material delivery savings.
- Establishes an internal case for potential enterprise licensing if value is proven.

4. Talent & Brand Advantage

- Builds consultant visibility in “next-gen delivery” initiatives — a differentiator in promotion, recruiting, and client perception.
- Showcases Slalom’s commitment to staying ahead of the curve on AI-driven delivery models.

Impact on the OpenAI Partnership

1. Co-Development of Enterprise Use Cases

- OpenAI gains structured feedback on Codex performance, user experience, and enterprise integration needs — data they rarely get from consultants.
- Slalom becomes a **trusted design partner**, bridging the gap between model capability and real-world delivery workflows.

2. Strategic Positioning for Future GTM

- Positions Slalom as a **preferred OpenAI implementation partner** for enterprise clients exploring AI-assisted engineering, code intelligence, or agentic systems.
- Creates pathways for **joint GTM offerings** — e.g., AI-assisted engineering workshops, AI code governance frameworks, or “AI-Augmented Delivery” accelerators.

3. Revenue & Influence Potential

- If the pilot demonstrates ROI, Slalom may advocate for enterprise Codex adoption internally and across clients — directly expanding OpenAI’s commercial footprint.
- Strengthens Slalom’s voice in shaping OpenAI’s **enterprise roadmap** (security, compliance, fine-tuning features, IDE integrations).

4. Trust & Responsible AI Governance

- Reinforces mutual credibility: OpenAI as a responsible tech provider; Slalom as a trusted system integrator ensuring safe and ethical AI deployment.

Summary in One Line

The Codex pilot makes Slalom a hands-on design partner to OpenAI — generating actionable insights on AI-assisted development, accelerating internal delivery innovation, and deepening a strategic alliance that positions Slalom at the forefront of enterprise GenAI enablement.

Aprajita

Aprajita - PreSales Talk

RFPs Details

Account extension, expansion are done to generate a healthy pipeline.

There are 2 parts to it - current and next fiscal year target, all the leaders do projection who is generating how much pipeline

Pipeline - leads -

While u r in the project, teher are more requirements coming up - it could be new services, or there is legacy system, fragmented data sources, they want to setup a new warehouse.

Business drives these requirements, because they do not get the information that they want.(my KPIs do not answer). The IT people generate projects - so whatever leads are with customer (e.g. salesforce, the customer says that i have these challenges in IT, how can slalom help), i will go i will create a lead, i will talk to practice and open up a oppty in salesforce, that is what is called a pipeline.

Multiple phases -

- **New logo** - not the existing customer,
- You are doing an extension in Existing account , no expansion (not as valued as expansion)
- **Extension + Expansion** : In your existing account, you are extending the piece of work + you are generating new scope of work which is getting generated from the existing work or new BU.
- RFPs - Either from existing customers, it can be from a new **logo** or new customer. If the project becomes bigger - 100 million + or 50 million + etc. Then the procurement team tells them to do RFP, do not give it to one single vendor. Its an internal process

As a solution architect , my role in pre-sales cycle:

While leaders are generating pipeline, you are brought in for solutioning. There are two phases to it:

1. **Initially top down** - we have some points on business pain points, technical pain points. We will tell you at high level point, and we provide a capability kind of view to it with little customization
2. **Bottoms Up** - it happens when leaders are aligned to a top down solution. At stage 1, 2 or 3, we ask the customer whether they like the solution, in that case, **then solutions architect and delivery people are involved. Besides solutioning, you do estimation and timelines and how many people are involved, what skillset is needed.** That

creates pricing and timeline. When you have to give the actual proposal then you do solutions architecting, where you want solutions architect to say - you know the problem statement and you can use these AWS services, so it's templatised as well as personalized based on the data points given.

Timeline - staggered approach, not everything at a start and end

Actual vs Forecasting - When you are in delivery, you create a timeline to estimate, how much of governance you need, hands-on, SME

T&M and Fixed Price - no need to get the details

RFP Cycle

1st phase - **respond to RFP**, if they like you they invite you

2nd Phase - You go for **Orals (optional)** - where you define and actually present the solution (client does 3-4 orals in a day - approx 2 hrs long), (people - CE client executive, practice lead (e.g. Salesforce - + data lead, pursuit lead, enterprise architect, SME for technologyclose to 3-4 people), and then they ask you questions and then get satisfaction

3rd Phase - If the project is too big, they do a **due diligence** workshop with you where they get into details and specifics, based on the facts from this workshop, you can refine your solution and present the final proposal.

Then the procurement and MSA happen and legal stuff.

—

Impact/Results of RFP, pre-sales work:

As a part of the solutions team, / solutions architect of D&A practice, I have contributed to 2 million amounts of the D&A pipeline working on multiple pre-sales deals and accounts.

Position of soln architect SME

I was part of the pre-sales account, where I was involved in responding to a proposal from the technical perspective where I have impacted close to 2-3 million dollars.

Multiple ways to get that target:

As a part of the delivery, as a project/pursuit lead, the team under you in that project, some factor accounts to your target.

If you have responded to multiple RFPs, and you have won the RFP, and you are part of that team, you will get credit that you were able to sell a certain amount.

Project yourself

My pitch:

Being part of the D&A practice one of the senior solutions architect who knows AWS very well, I was part of multiple proposals, which involves certain...i was creating technical solution documents, engaged in doing workshops with the client to try and understand, collect the data point to do initial assessment, to do initial discovery to make sure my solution aligns with the long term delivery plan. In that phase i have impacted 2-3 million dollar of proposals through RFPs.

How did you do it - i worked with the pod, i worked with the account team, i worked with RFP timelines, depending on what their requirements is, we responded in a certain way,

i have also helped with the SOW, writing down the activities, what deliverables, how the RACI will be defined, what are the technical components or automation, which should be involved in the solution.

Impact - it generates the pipelines for our D&A,bring business to company, helping Slalom grow.
https://en.wikipedia.org/wiki/Responsibility_assignment_matrix

Partners

Alliance Partners -

Sometimes we have do POC or discovey with the client and we need to be funded. Since we do not want to charge the client, we reach ut to our technology partners to get funding from them. Our technology partners fund us when we are proposing a good workload, may be aws or GCP or something, what services we are proposing, what is the data volume, what integrations are involved, how much data volume is expected in future etc.

Basically aws, gcp etc, try to understand how muhc we are helping to increase their business. There is some benchmark to it (for small workload - x amount, bigger workload - y amount).

My pitch:

I also have experience working with our alliance partners within Slalom, in which I have worked to get some MAP funding, or **funding** from technology partners, which we have invested in doing small POCs, or trying out use cases at the client side.

You also worked with multiple practices. I have worked with Salesforce practice, to talk about the **salesforce data cloud integration**,

You should show - you can coordinate with multiple things, you can pull multiple threads to get a proposal done.

Funding Details:

For partners to do the business, they can do it in 2 ways

Directly go to the customer - challenge is they can buy the solution but who will implement it, that is where consultant companies come in. Slalom partners

Partners want us to propose our solution, do POCs, we will fund you to try out our capabilities. You know how to implement, you know the pain points of the customer, you sell to me, and we will sell to you.

Sometimes the clients go for projects to aws, then aws will say we will give you so much discount, if you work with Slalom. That's the role of alliance partners, to create that relationship.

As a solns architect:

You reach out to multiple alliance partners, to vet out a solution from different aspects - 1. Technical aspect to compare solns. 2. From a licensing perspective what is better from a cost perspective.

(Licensing - services use .. cost etc, subscription cost - aws, gcp, vs azure discount. For hyperscale - multicloud, the license cost are mostly same, the main difference that comes is license ...how much slalom can get discount for us.)

As a solns architect, I need to make sure that with the given data points, at a high level its doable, feasible, scalable for the future, and sustainable.

Also you need to think end-to-end , which apps are integrating into cloud, in multi cloud, when two systems talk to each other, sap or aws charges its own. At an enterprise level, you are looking into integration perspective, and costing perspective,

So what - if you partner with alliance partner, your motive is to get money from them, and you give

Price to win - we maintain a margin to project, 20% , 30% if the project is not finishing up on time, then margin helps.

We are getting ur poc or discovery free as we are getting partner funding. With that partner funding slalom are paid. With the commitment that you use my solution and going to get business for long terms.

Dont take client name -

Salesforce - big CRM company

Apple - fruitstand

It maybe aws is - Banking sector company gets more discount vs small company, it depends on what band they have put the company in.

Gold partner - base don how much business has sloam done to aws...

EXamples::

extension/expansion to generate the lead

What kind of RFPs you worked on (example of RFPs)

Have a fact in your mind

Work with all whom - projects , RFPs, client interaction - any opportunity generated

=====

I just wanted to share an achievement that I'm proud of...

11:02

I recently completed the data discovery and recommendations phase and presented my findings to the customer. The feedback was very positive! The primary contact on the customer side even mentioned that this was the first time in their tenure at Salesforce that they've been able to put together this type of documentation, given how complex and challenging the ecosystem is.

11:04

Also, it looks like we might be able to secure more people for the project. Check out this message:

Hongwu Wang

5:15 PM

Hi @Madhurima Saxena Thank you for walking us through the recommendations on our data. We just kicked off a project on revamping our data environment in preparation for the future spiff data. I wonder if someone from the team can help us PM the project, as well as acting as the consultant? Also, I will set up a separate meeting to learn more about the team.

AWS MAP

The AWS Migration Acceleration Program (MAP) is a program offered by Amazon Web Services (AWS) designed to help businesses accelerate their journey to the cloud. It provides a

comprehensive set of services, best practices, tools, and financial incentives to facilitate migration projects. Here are some key aspects of the AWS MAP program:

Assessment Phase: This initial phase helps organizations evaluate their readiness for cloud migration. It involves assessing current IT environments, identifying the business case for migration, and creating a detailed migration plan. AWS and its partners provide tools and resources to help with cost modeling and identifying potential savings.

Mobilization Phase: During this phase, organizations focus on building the necessary foundation for a successful migration. This includes setting up a Cloud Center of Excellence (CCoE), developing operational skills, and addressing any gaps in security, compliance, and governance.

Migration and Modernization Phase: The final phase involves executing the migration plan, which includes moving workloads to the AWS cloud. AWS and its partners provide tools and services to automate and accelerate the process. This phase also includes modernizing applications to take full advantage of cloud-native features and services.

Training and Support: AWS MAP offers training resources to help teams build cloud skills and knowledge. AWS and its partners provide ongoing support throughout the migration process to ensure a smooth transition.

Financial Incentives: To offset migration costs, AWS MAP provides financial incentives, which can include AWS service credits and funding for partner-led services.

Partner Network: AWS collaborates with a network of partners who are experienced in cloud migrations. These partners can provide additional expertise and support during the migration process.

The AWS MAP program is particularly beneficial for organizations looking to migrate large-scale workloads or complex IT environments to the cloud. It's designed to reduce risk, enhance efficiency, and accelerate the time to realize the benefits of cloud adoption. If you are considering utilizing the AWS MAP program, it is recommended to work with an AWS representative or an AWS Partner to tailor the program to your specific needs and goals.

stripe

The expectation is that you're working with them on deals that are in flight that they have where they need stripe help in the deal.

or that you're helping them build solutions or white papers or whatever it might be because a lot of white people's at AI right now. I see that's like we might be contributing to their AI strategy from a payments and kind of like money movement perspective. So getting in and plugging in and working with them, and that's 2 or 3 things, right, that's one deep knowledge of stripe, 2, that's deep knowledge or deep relationships with people at Slalom, and then 3, that is, uh, like an actual target test case use case like want be an industry, it might be a specific customer, it might be like there has to be someone who's actually going to be the target for this. And having enough knowledge to know who that is, to go like, OK, slalom are in this place, they're going to do some AI agentic commerce thing or whatever it is.

I know they're really good at automotive, so I'm going to make sure that on my side of the stripe end I've got whoever's good at automotive to help with this, but I'm the lead architect on what that might look and feel like. So it's not like a day to day thing, it's more just like quarter over quarter, we're constantly working on that type of stuff.

And then there's also like slightly more reactionary, like, there's deals coming in from those partners, we're working with them on that, that might mean like helping them with the demos, helping them with their presentations, whatever it might be, um, and then the other one is making sure that we're at the forefront of enabling those partners to have stripe resources that are certified, that are good, that are capable, that are able to like bring scale to stripe, because like we're not gonna, we're not gonna fill buildings around the world with people who go to market for us or do implementation for us, we're going to have to leverage partners for that. So that's kind of like

our job is to be the technical like field CTO that is out there with those partners making sure that they're all capable of doing that, so that's why we were measured on those white pipeline and all those kind of things are kind of standard operating metrics because that's the best like lagging indicator that's the stuff that we're doing this today is working.

Does that, does that answer that? Does that make sense?

Links

Links

<https://www.kdnuggets.com/top-sql-patterns-from-faang-data-science-interviews-with-code>

<https://www.databricks.com/company/careers/interview-prep>

Important Notes

[Rejected Again? Here's the Truth Nobody Tells You About Hiring | by George J. Ziogas | Psychology of Workplaces](#)

<https://skillroads.com/free-resume-review>

Databricks IW

1. Know about Databricks - Lunch and Learn is important. Compared with Snowflake.
2. Do DBx professional
3. References/Links:
 - a. <https://www.youtube.com/watch?v=hahPTeXJEBk>
 - b.  How to Ace the Databricks Solutions Architect Technical Interview | Tips fro...
 - c. https://www.glassdoor.com/Interview/Databricks-Solutions-Architect-Interview-Questions-EI_IE954734.0,10_KO11,30_IP2.htm?filter.jobTitleFTS=Solutions+Architect
 - d. https://www.reddit.com/r/salesengineers/comments/1iawyr3/databricks_solution_architect_interview/
 - e.
 - f. https://www.reddit.com/r/salesengineers/comments/1iawyr3/databricks_solution_architect_interview/
 - g. <https://twodegrees1.sharepoint.com/:p:/t/TNCRFP/EZMNpZ0OYjxCuZ7KEzAffzEB0f4pGImsAVgjMO7WyrN5rw?e=NsM8wd&ovuser=9ca75128-a244-4596-877b-f24828e476e2%2Cmadhurima.saxena%40slalom.com&clickparams=eyJBcHBOYW1lIjojVGh0bXN0bGVza3RvcCIsIkFwcFZlcnNpb24iOiI1MC8yNTEwMDIxMzgxMSJ9>
 - h.
4. Take notes at
 - a. Jeffs intro - been with slalom - 12:57 min

<https://www.youtube.com/watch?v=Tfpdwf45eK8>

https://youtu.be/iHmCysEsp5o?si=bqr72ubCspj_bpAP

MS

“What’s your experience with Azure?”

Use this **script**:

“My hands-on experience has primarily been with AWS and Databricks, but the core of my work has been cloud-agnostic data architecture — designing lakehouse platforms, modernizing legacy warehouses, enabling analytics and AI, and **guiding customers through adoption**.

I’ve worked closely with Azure-based teams and solutions, and I’m very comfortable transferring these patterns to Azure. What’s important to me is choosing the right architecture for the customer, and I ramp very quickly on cloud-specific services.”

Why Microsoft?

I want to work at Microsoft because of the impact — Microsoft is shaping how enterprises responsibly adopt data and AI at global scale. This role excites me because it combines customer-facing architecture, data modernization, and business impact. It’s where I do my best work: helping customers design secure, scalable, AI-ready platforms and guiding them from idea to production and adoption.

Azure’s competitive advantage

Azure’s competitive advantage really comes down to how well it serves enterprise customers end-to-end.

First, Azure has a very strong enterprise and hybrid story. Many organizations already run Microsoft across identity, productivity, and core systems, and Azure integrates naturally into that environment. That lowers migration risk and makes hybrid and regulated workloads much easier to support compared to a cloud-only approach.

Second, Azure offers a unified data and AI platform. Customers can move from data ingestion and analytics to BI and AI within a single ecosystem — using services like Databricks, Fabric, Power BI, and Azure AI. That reduces tool sprawl and accelerates time to value, especially for AI-driven use cases.

Third, governance and security are built in from the start. Azure’s identity, access control, and compliance capabilities are enterprise-grade and deeply integrated, which is critical as customers start putting more sensitive data and AI workloads into production.

Finally, Azure benefits from the broader Microsoft ecosystem. Integration with tools people already use every day — like Microsoft 365 and Power BI — drives adoption and makes insights more actionable across the business.

Overall, Azure differentiates itself by combining enterprise trust, hybrid flexibility, and a cohesive data-to-AI platform that helps customers modernize responsibly and at scale.

- *How did you discuss the value of the solution, including ROI and TCO?*
- *What was your solution's value proposition in comparison to other competitor solutions?*

When discussing value, I always start by anchoring the conversation on the customer's business goals rather than the technology itself. In this case, the primary objectives were faster time to insight, scalability for future analytics and AI use cases, and reducing the operational overhead of managing fragmented data systems.

From an ROI perspective, we focused on both hard and soft returns. **On the hard side, we quantified cost savings from consolidating legacy systems**, reducing infrastructure and licensing spend, and lowering operational effort through managed services and automation. **On the soft side, we highlighted faster delivery of analytics**, improved data trust, and the ability for teams to self-serve insights — which directly impacted decision-making speed and productivity.

For TCO, we looked beyond just compute costs. We evaluated total ownership across infrastructure, operations, maintenance, and future scalability. The proposed architecture reduced long-term costs by simplifying the platform, minimizing manual data pipelines, and allowing the customer to scale usage up or down based on demand rather than over-provisioning.

When comparing against alternative solutions, our value proposition was flexibility and future readiness. Competing options either locked the customer into a more rigid architecture or required significant custom engineering to support analytics and AI at scale. Our solution provided an open, scalable data foundation that could evolve as the customer's needs grew — without a major re-architecture.

Ultimately, the decision came down to balancing cost, risk, and long-term value. The customer chose our approach because it delivered near-term ROI while also positioning them for future innovation, particularly around advanced analytics and AI.

AWS Data Engineering

AWS Data Engineering
has 4 levels to it:

– Level 1: Ingesting & Storing Data

Start by learning the foundations of AWS data services:

- S3 for data lakes (folders, prefixes, lifecycle policies)
- AWS Glue Crawlers + Data Catalog for schema discovery
- Kinesis / AWS DMS for streaming + CDC ingestion
- IAM basics (roles, policies, S3 bucket access)

With just these basics, you can already build working ETL pipelines.

– Level 2: Transforming & Querying Data

Move from storing data to making it usable:

- Glue ETL & PySpark jobs (batch transformation)
- AWS Lambda for lightweight event-driven processing
- Athena + S3 for serverless SQL on data lakes
- Redshift for warehousing and complex analytics

This is where your data becomes queryable and structured for analytics

– Level 3: Building Scalable Data Platforms

Upgrade from pipelines to full data platforms:

- Lakehouse architecture with Iceberg/Delta on S3
- Glue Workflow/Step Functions for orchestration
- Data partitioning, file formats (Parquet, ORC, Avro)
- Performance tuning (compaction, distribution keys, sort keys)

Here's where you shift from "data exists" to "data is optimized and reliable."

– Level 4: Operating at Scale

Finally, learn to make your platforms efficient, secure, and enterprise-ready:

- EMR + Spark clusters for high-volume processing
- Data quality + observability (Great Expectations, Deequ, CloudWatch)
- Cost optimization (S3 tiers, Redshift RA3, Glue job tuning)
- Security & compliance (KMS encryption, VPC endpoints, Lake Formation, GDPR/SOC2)
- Streaming at scale with Kinesis Data Streams / Firehose / MSK

ML

1. Explain the ML lifecycle on Databricks in under 60 seconds.

30-sec answer:

“On Databricks the ML lifecycle is unified across data, features, training, and deployment.

We ingest raw data into Delta Lake so it's versioned and reproducible.

Feature engineering happens in notebooks, pipelines, or Delta Live Tables and those features are stored in the Feature Store.

MLflow tracks experiments, models, and parameters.

When a model is ready, it goes into the **Model Registry** where we manage staging → production promotion with CI/CD.

Finally, we serve it through **Model Serving** or batch jobs, and use Unity Catalog & monitoring to track drift, quality, and lineage end-to-end.”

3. How would you scale training on TBs of data?

30-sec answer:

“I use Databricks' distributed compute — Spark for data prep, and MLflow + distributed training libraries like Horovod or PyTorch DDP for model training.

Because the data lives in Delta Lake, we can optimize with partitioning, Z-Ordering, and caching, then run parallel hyperparameter tuning with Spark-based trials.

So instead of trying to train on one machine, the workload scales horizontally across the cluster.”

4. Batch scoring vs online scoring — when and how?

30-sec answer:

“Batch scoring is best when predictions are needed periodically at scale — like nightly risk scores.

It runs as a scheduled job reading from Delta Lake and writing predictions back.

Online scoring is for low-latency use cases like recommendations or fraud.

Databricks Model Serving exposes the model as a REST endpoint with auto-scaling.

Both use the same model in MLflow Registry to keep governance consistent.”

5. How does MLflow help enterprise ML teams?

30-sec answer:

“MLflow standardizes experiments, model packaging, and lifecycle governance.

Teams can track metrics, compare runs, reproduce results, and version models in the Registry.

It becomes the central nervous system for ML because data scientists, engineers, and SREs all operate from the same model registry and lineage — which avoids shadow models, manual handoffs, and deployment inconsistencies.”

6. Explain point-in-time correctness in the Feature Store.

30-sec answer:

“It means features are computed using only the data that was available at that exact timestamp.

When you build a training set, the Feature Store performs a time-aware join so future information cannot leak into the model.

This solves one of the biggest problems in enterprise ML — reproducible, leakage-free datasets.”

7. How do you evaluate a RAG pipeline?

30-sec answer:

“I evaluate RAG on three levels:

1. **Retrieval metrics:** precision@k, recall@k, coverage.

2. **Generation metrics:** faithfulness, hallucination rate, groundedness.

3. **System metrics:** latency, cost, token usage.

On Databricks, Lakehouse Monitoring tracks these automatically and logs them to MLflow so teams can compare versions and detect drift.”

8. When should a customer NOT use deep learning?

30-sec answer:

“When the dataset is small, tabular, or when interpretability is more important than raw accuracy.

Tree-based models like XGBoost often outperform deep learning on structured data with less cost and complexity.

Deep learning makes sense for unstructured modalities or when scale gives a real advantage — otherwise, it’s overkill.”

9. How do you handle schema drift in ML pipelines?

30-sec answer:

“I design pipelines to be schema-aware.

Delta Lake’s schema evolution and expectations can validate incoming data.

Feature engineering code should be defensive — missing columns default to null or fallback logic.

We track schema versions in MLflow so when a new schema appears, we retrain or branch the model.

This avoids silent degradation.”

10. How do you operationalize a model with CI/CD?

30-sec answer:

“Code, features, and models all live in Unity Catalog.

A Git-triggered workflow runs tests, builds features, trains the model, and registers it in MLflow.

Promotion from Staging → Production is gated through approval or automated tests.

Production models are served through Model Serving or batch jobs, and monitored for drift using Lakehouse Monitoring.”

11. Explain difference between feature engineering at inference vs training.

30-sec answer:

“At training time you compute features over historical data — often aggregated or windowed.

At inference time you only use data available *right now*.

The Feature Store ensures both use the same code path so features stay consistent.

If these diverge even slightly, you get poor real-world accuracy.”

12. What would you do if a model’s performance dropped unexpectedly?

30-sec answer:

“First check data drift, schema changes, or missing features — those cause most drops.

Then compare current inference features with training features using MLflow logs.

Next check model drift using statistical monitoring.

If the issue is real, we retrain using the latest Delta snapshots and promote a new version through the Registry.”

✓ Part 2 — A full mock customer scenario where you design an ML platform end-to-end

Here's a polished, interview-ready scenario that makes you sound like a **strong Databricks SA**.

★ Mock Customer Scenario (Databricks SA Case Study)

Customer:

RetailCo, a global retail chain with 2,000 stores.

Goal:

They want to build a **real-time recommendation and forecasting platform** to improve:

- product recommendations
- demand forecasting
- store workforce planning
- supply-chain optimization

Their pain points:

- Data is split across Snowflake, S3, and on-prem SQL Server
 - DS teams manually build training datasets → no reproducibility
 - Production ML is ad-hoc: notebooks emailed around, no CI/CD
 - They want to explore GenAI for store-manager Q&A
-

★ Your Proposed Architecture (End-to-End)

1. Data Ingestion & Storage

- Ingest structured & unstructured data using Auto Loader into **Delta Lake** on Databricks.
- Use DLT (Delta Live Tables) for **quality checks**, expectations, and schema inference.
- All data is governed through **Unity Catalog**.

2. Feature Engineering & Feature Store

- Build reusable features:
 - customer embeddings
 - RFM metrics
 - transaction sequences
 - store-level aggregations
- Store all features in the **Databricks Feature Store**.
 - ensures point-in-time correctness
 - eliminates leakage
 - shared across DS teams

3. Model Training

- Use **MLflow** to track all experiments.
- Distributed training using:
 - **XGBoost** for forecasting

- **Deep Learning** (PyTorch Lightning) for recommendations
- Hyperopt or SparkTrials for scalable hyperparameter tuning.
- Save the best model to the **MLflow Model Registry**.

4. Deployment

Two paths:

Batch Scoring

- Nightly forecasts written back into Delta.
- BI tools (Tableau/PowerBI) read from Delta for dashboards.

Real-Time Serving

- Deploy recommendation model via **Databricks Model Serving**.
- Low-latency endpoints backed by autoscaling clusters.

5. Monitoring & Governance

- Use **Lakehouse Monitoring** for:
 - data drift
 - model drift
 - feature skew
 - latency & cost metrics
- Alerts back to Slack or PagerDuty.

6. GenAI Add-On (Execs love this)

RetailCo wants a **store-manager assistant** that answers questions like:

“Why is forecast down for Store #812?”

Proposed RAG pipeline:

- Documents stored in Delta (PDFs, SOPs, plans)
- Chunk → embed → store embeddings in **Databricks Vector Search**

- Use **DBRX** or Llama-3 fine-tuned using Mosaic AI
 - Serve through Model Serving with built-in monitoring
 - Track hallucination rate + groundedness with MLflow
-

★ 30-Second Summary You Could Say in the Interview

“We unified RetailCo’s ML lifecycle on Databricks. All raw data lands in Delta Lake through Auto Loader. We built reusable features in the Feature Store with point-in-time correctness, then trained models for forecasting and recommendations with MLflow.

The models move through Model Registry into batch pipelines or real-time Model Serving.

Everything — data, features, models, lineage — sits under Unity Catalog so governance is automatic.

We added a RAG assistant using Vector Search + DBRX so store managers get contextual answers.

This gives RetailCo an end-to-end, governed, scalable ML platform.”

If you want, I can also prep:

- ✓ 30-second “Why Databricks for ML?” pitch
- ✓ 10 deep-dive questions *you should ask them*
- ✓ A second mock scenario for a healthcare or financial customer
- ✓ A version of the answers in your *storytelling / conversational* style

Phrases to use

Phrases to use

1. Delta's versioning also guarantees reproducibility so we can rebuild the exact training dataset if needed."
2. Databricks Model Serving exposes the model as a REST endpoint with auto-scaling.
3. DS teams manually build training datasets → no reproducibility
4. Everything — data, features, models, lineage — sits under Unity Catalog so governance is automatic.
5. Data is split across Snowflake, S3, and on-prem SQL Server
- 6.